

# Assignments week 5

Son Cao

2025-10-01

## Assignment 1 - Removing values from a binary heap

My code:

```
#include <iostream>
#include "utils.h"
#include <vector>

//min-heap

using namespace std;

void BubbleUp(vector<int>& heap, int i){
    while(i > 0){
        int parent = (i - 1) / 2;
        if(heap[i] < heap[parent]){
            swap(heap[i], heap[parent]);
            i = parent;
        } else break;
    }
}

void BubbleDown(vector<int>& heap, int i){
    int n = heap.size();
    while(true){
        int left = 2 * i + 1;
        int right = 2 * i + 2;
        int smallest = i;

        if(left < n && heap[left] < heap[smallest]) smallest = left;
        if(right < n && heap[right] < heap[smallest]) smallest = right;

        if(smallest != i){
```

```

        swap(heap[i], heap[smallest]);
        i = smallest;
    } else break;
}
}

int main() {
/* TODO:
   Write a program that reads a binary *min-heap* from its standard input, given as an
   and an integer that represents a value to be removed from the heap.

   The program must remove the specified value from the heap while maintaining the heap
   property, and print the resulting heap in the same format as the input. If the value
   is not present in the heap, the program must print the original heap unchanged.

   Example input:
   [3, 4, 10, 23, 50, 90] 50
   Example output:
   [3, 4, 10, 23, 90]
*/
vector<int> heap;
int value;
cin >> heap >> value;

int idx = -1;
for(int i = 0; i < (int)heap.size(); i++){           //O(n)
    if(heap[i] == value){ //find index of value
        idx = i;
        break;
    }
}

if(idx != -1){
    //replace with last element
    heap[idx] = heap.back();
    heap.pop_back(); //O(1)

    if(!heap.empty()){
        if(idx > 0 && heap[idx] < heap[(idx - 1) / 2]){
            BubbleUp(heap, idx); //O(log n)
        } else {
            BubbleDown(heap, idx);
        }
    }
}

cout << heap << endl;
}

```

```

    return 0;
}

```

Even though restoring the heap takes only  $O(\log n)$ , the dominant step is the linear search to find the element => that's why the overall complexity is  $O(n)$ .

Time complexity: this algorithm has a time complexity of  $O(n)$ , because the algorithm must linearly search for the value to remove ( $O(n)$ ), and then restore the heap property, which takes  $O(\log n)$ . The linear search dominates the total time.

## Assignment 2 - Heapify

My code:

```

#include <iostream>
#include "utils.h" // for reading vectors
#include <vector>

using namespace std;

int BubbleDown(vector<int>& heap, int i){
    int swaps = 0;
    int n = heap.size();
    while (true){
        int left = 2 * i + 1;
        int right = 2 * i + 2;
        int smallest = i;

        if(left < n && heap[left] < heap[smallest]) smallest = left;
        if(right < n && heap[right] < heap[smallest]) smallest = right;

        if(smallest != i){
            swap(heap[i], heap[smallest]);
            swaps++;
            i = smallest;
        } else break;
    }
    return swaps;
}

int heapify(vector<int>& arr){
    int swaps = 0;
    for(int i = arr.size()/2 - 1; i >= 0; --i){ //the last non-leaf node and goes up to the
        swaps += BubbleDown(arr, i);
    }
}

```

```

    return swaps;
}

//make_heap(arr.begin(), arr.end(), greater<int>()); doesnt return swaps

int main() {
    /* TODO:
        Write a program that reads an array from its standard input, given as an array between square brackets (e.g. `[1,3,5,7,9]` - see week 1), and prints the number of *swaps* would be performed when turning this array into a valid binary *min-heap* using the `std::make_heap` function from the C++ Standard Library.
    */
}

```

*You can't use the `std::make\_heap` function directly, instead you must implement the heapify algorithm as described in the lecture, and count the number of swaps performed during the process.*

*The time complexity of your solution must be  $O(n)$ , where  $n$  is the number of elements in the input array.*

```

Example input:
[1, 5, 6, 2, 3, 4, 7]
output:
2
*/
vector<int> arr;
cin >> arr;

int swaps = heapify(arr);
cout << swaps << endl;
return 0;
}

```

Time complexity: this algorithm has a time complexity of  $O(n)$ , because it builds the heap from the bottom up. Most elements are near the bottom of the tree and need only a few swaps to fix their position, so the total amount of work adds up to a value that grows roughly in proportion to the number of elements.

## Assignment 3 - Running median

My code:

```

#include <iostream>
#include <queue>
#include <iomanip>

using namespace std;

```

```

int main() {
    /* TODO:
        Write a program that reads a sequence of numbers from its standard input and, after
        prints the median of all numbers read so far.

        You should use two binary heaps to implement the required functionality: a max-heap
        the lower half of the numbers, and a min-heap to store the upper half of the numbers.

        The time complexity of your program must be O(n log n), where n is the total number
    */
    priority_queue<int> lowers; //priority_queue<int> is a max-heap by default => keep the
    priority_queue<int, vector<int>, greater<int>> uppers; //min-heap => keep the upper half

    double median;
    int x;

    while(cin >> x){
        if(lowers.empty() || x <= lowers.top()){
            lowers.push(x);
        } else {
            uppers.push(x);
        }

        if(lowers.size() > uppers.size() + 1){
            uppers.push(lowers.top());
            lowers.pop();
        } else if (uppers.size() > lowers.size() + 1){
            lowers.push(uppers.top());
            uppers.pop();
        }

        if(lowers.size() == uppers.size()){
            median = (lowers.top() + uppers.top()) / 2.0;
        } else if (lowers.size() > uppers.size()){
            median = lowers.top();
        } else {
            median = uppers.top();
        }
        cout << median << " "; //cannot do cout << median << endl; cuz it prints ex: 1\n2\n
    }
    cout << endl;
    return 0;
}

```

Time complexity: this algorithm has a time complexity of  $O(n \log n)$ , because

for each of the n numbers, insertion and balancing operations in a heap take  $O(\log n)$  time (because inserting into a heap requires rearranging elements to keep the heap order). The constant-time median calculation doesn't affect the overall complexity.

## Assignment 4 - Ordering coordinates

My code:

```
#include <iostream>
#include <unordered_map>
#include <vector>
#include <algorithm>

using namespace std;

struct Coord{
    int x, y;
    bool operator==(const Coord& other) const{
        return x == other.x && y == other.y;
    }
};

struct CoordHash{      //converts a key (like a coordinate) into an integer number that can be compared
    size_t operator()(const Coord& c) const{           // ^ (XOR) -> mix the bits together
        return hash<int>()(c.x) ^ (hash<int>()(c.y) << 1); // << 1 -> shift one of them left
    }
};

int main() {
    /* TODO:
       Write a program that inputs a string consisting of the characters "U", "D", "L", and "R", representing movements Up, Down, Left, and Right on a 2D grid, starting from the origin (0, 0).
       The program must then compute and print the top three of *the most visited unique coordinates* that are visited during the sequence of movements, including the starting point (0, 0).
       Ties should be broken by the order in which the coordinates were first visited.
       For example, given the input string "UUDLURD", the program must print `(0, 1), (0, 2), (-1, 1)` because the visited coordinates are (0, 0), (0, 1), (0, 2), (0, 1), (-1, 1), (-1, 2), (0, 2).
       The time complexity of your solution must be  $O(n + 3 \log(n))$ , where n is the length of the input string.
    */
}
```

To efficiently store and check for unique coordinates, you should use a hash table (``std::unordered_set`` or ``std::unordered_map`` in C++). Note that you will have to define a comparison function for the coordinate pairs, as well as an equality operator. Also, you must convert the coordinates into a vector or similar structure to be able to sort them by their visit counts.

```

As a final note, you are not allowed to use the nth_element algorithm from the standard library to sort the coordinates. Instead, you must implement your own sorting logic.

/*
string moves;
cin >> moves;
unordered_map<Coord, pair<int, int>, CoordHash> visits;
//Coord => visits[{x, y}]
//pair<int, int> => ex: {1, step}
//first (1) = number of times this coordinate was visited
//second (step) = step index when this coordinate was first visited (used for tie-breaker)

int x = 0, y = 0;
int step = 0;
visits[{x, y}] = {1, step++}; //ex: visits[{0,1}] = {3, 1}; => Coordinate (0,1) was visited once at step 1

//the moves
for(char move : moves){
    if(move == 'U') y++;
    else if(move == 'D') y--;
    else if (move == 'L') x--;
    else if(move == 'R') x++;

    Coord current = {x, y}; //current will be used as a key in the unordered_map visits
    if(visits.find(current) == visits.end()){ //visits.find(current) looks for current coordinate
                                                //visits.end() represents "not found"
        visits[current] = {1, step++};
    } else {
        visits[current].first++; //We do not change the first-visit step (second), because we
                                //are only interested in the count of visits
    }
}

//sort to get top 3 //unordered_map does not store items in order, so we need to copy all entries
vector<pair<Coord, pair<int,int>>> all(visits.begin(), visits.end()); //copy all entries
// pair<Coord, pair<int,int>>
// all[i].first => the coordinate
// all[i].second.first => number of visits
// all[i].second.second => first-visit step

sort(all.begin(), all.end(), [] (auto &a, auto &b){
    if((a.second.first != b.second.first))
        return a.second.first > b.second.first; //true if a goes before b, false if a goes after
    return a.second.second < b.second.second; //if visit counts differ => bigger count
})

```

```

}); // if counts are equal => smaller first

int count = 0;
for(auto &p : all){ // p: reference to each element
    cout << "(" << p.first.x << ", " << p.first.y << ")";
    count++; //track how many coordinates we printed
    if(count == 3 || count == (int)all.size()) break;
    cout << ", ";
}
cout << endl;
return 0;
}

```

Time complexity: this algorithm has a time complexity of  $O(n + \log n)$ , because it processes each move once to update coordinates and counts using an unordered\_map, which gives  $O(1)$  average time per operation, so this part is  $O(n)$ . And it then sorts the coordinates to find the top 3 most visited ones. Sorting normally costs  $O(k \log k)$  (where  $k$  is the number of unique coordinates), but since we only need the top 3, the cost simplifies to  $O(3 \log n) \approx O(\log n)$ .

## Assignment 5 - Linear probing

My code:

```

#include <iostream>
#include "utils.h"
#include <vector>
#include <string>
#include <bits/hash_bytes.h>

using namespace std;

struct string_hash{
    size_t operator()(const string& str) const{
        return std::Hash_bytes(str.data(), str.size(), static_cast<size_t>(0xc70f6907UL));
    }
};

int main() {
    /* TODO:
       Write a program that inputs an integer followed by a list of strings (given as a comma-separated list between square brackets, e.g. [apple, banana, cherry], use the utils.h header to read the input)
       The list represents the strings to be inserted into a hash table that uses linear probing resolution, and the integer represents the number of slots in the hash table.
    */
}

```

*The program must insert the strings into the hash table in the order they are given resulting hash table as a comma-separated list between square brackets, where empty represented by the string "EMPTY".*

*The hash table must implement a \*set\*, meaning that duplicate strings should not be times. In other words, if the same string appears more than once in the input list, appear once in the output hash table.*

*You must use the `std::hash<std::string>` struct from the C++ Standard Library to compute of a string. When determining whether a string is already present in the hash table, search the entire table. Instead, you should use the same probing sequence that would be used for linear probing, stopping when you either find the string or reach an empty slot.*

*Example input:*

*((2 3 4) 5 (6 7 8)) (3 5 10)*

*Example output:*

*[3, 5]*

*\*/*

```
// Hint: write a function to find an empty slot using linear probing
int table_size;
char c;
cin >> table_size >> c;
vector<string> table(table_size, "EMPTY");
string word;

string_hash hasher;

while(cin >> word){
    if(word.back() == ',') word.pop_back(); //check for "[" first
    if(word.back() == ']') word.pop_back();

    size_t index = hasher(word) % table_size; //tryna find index in a table (size_t in

    for(int i = 0; i < table_size; i++){
        size_t probe = (index + i) % table_size; //wrap around the table check for emp

        if(table[probe] == "EMPTY"){
            table[probe] = word;
            break;
        } else if (table[probe] == word){
            //if duplicate => ignore
            break;
        }
    }
    if(word.back() == ']') break;
}
```

```

    }
    cout << "[";
    for(int i = 0; i < table_size; i++){
        cout << table[i];
        if(i != table_size - 1) cout << ", ";
    }
    cout << "]" << endl;
    return 0;
}

```

For each string in the input, the program first calculates its hash, which takes about  $O(k)$  time where  $k$  is the length of the string. Then it tries to insert the string into the hash table using linear probing. In the worst case, if the table is nearly full or there are lots of collisions, it might have to check every slot in the table, which is  $O(m)$  where  $m$  is the table size.

Time complexity: this algorithm has a time complexity of  $O(n * (m + k)) = O(n * m + n * k)$

## Assignment 6 - Separate chaining, rehashed

My code:

```

#include <iostream>
#include "utils.h"
#include <vector>
#include <forward_list>
#include <bits/hash_bytes.h>

using namespace std;

struct string_hash{
    size_t operator()(const string& str) const {
        return std::hash_bytes(str.data(), str.size(), static_cast<size_t>(0xc70f6907UL));
    }
};

//insert with separate chaining(prepending)
void insert_string(vector<forward_list<string>>& table, const string& word, string_hash& hasher){
    size_t index = hasher(word) % table.size();

    //check for duplicate
    for(auto& s : table[index]){
        if(s == word) return;
    }
    table[index].push_front(word); //prepending
}

```

```

}

int main()
{
    /* TODO:
        Write a program that reads a integer followed by a list of strings from its standard input.
        The list represents the strings to be inserted into a hash table that uses *separate chaining* for collision resolution, and the integer represents the number of slots in the hash table.
        The program must insert the strings into the hash table in the order they are given, resulting in a linked list of strings. The output should be a comma-separated list between square brackets, where each slot is represented as a linked list in the form `[<string1> -> <string2> -> ...]` (empty slots will be represented as a pair of brackets, i.e. "[]").
        The hash table must implement a *set*, meaning that duplicate strings should not be inserted more than once. In other words, if the same string appears more than once in the input list, it should appear once in the output hash table.
        To determine the index for a string, you *must* use the `std::hash<std::string>` standard library function provided in the Standard Library to compute the hash value of the string.
        When determining whether a string is already present in the hash table, you must not use the `find` function. Instead, you should use the same probing sequence that would be used for insertion, which is to continue probing until either you find the string or reach the end of the linked list.
    */

    // Hint: use a vector of singly linked lists as the underlying data structure for the hash table
    int k, n;
    vector<string> words;
    cin >> k >> words >> n;

    string_hash hasher;
    //build initial table
    vector<forward_list<string>> table(k);    //using forward_list because it's -> (forward pointer)
    for(auto& word : words){
        insert_string(table, word, hasher);
    }

    //rehash into new table
    vector<forward_list<string>> new_table(n);
    for(auto& bucket : table){
        for(auto& word : bucket){
            insert_string(new_table, word, hasher);
        }
    }
}

```

```
}

//print the final rehashed table
cout << new_table << endl;
return 0;
}
```