

# Assignments week 7

Son Cao

2025-10-16

## Assignment 1 - Topological sorting

My code:

```
#include <iostream>
#include "utils.h"    // for reading vectors
#include <vector>
#include <unordered_map>
#include <unordered_set>
#include <string>
#include <algorithm> //std::min

std::unordered_map<std::string, std::vector<std::pair<std::string, int>>> graph;
std::unordered_map<std::string, int> state;
std::vector<std::string> topo;
bool has_cycle = false;

void dfs(const std::string& u){
    state[u] = 1;
    for(auto &[v, w] : graph[u]){    //v: dest node(), w: edge weight(int)
        if(state[v] == 0) dfs(v);
        else if(state[v] == 1) has_cycle = true;
    }
    state[u] = 2;
    topo.push_back(u);    //C->B->A
}

int main() {
    /* TODO:
       Write a program that reads a list of edges representing a *weighted directed* graph
       The graph is given as a comma-separated list between square brackets, e.g. `[(A, B,
       Use the class `Edge<T>` defined in `utils.h` to represent the edges of the graph.
```

```

The program must check if the graph contains a cycle, and if it does not it must compute the shortest path from start to end.
It must do so by topologically sorting the nodes of the graph and then using this ordering to compute the shortest path.

*/
std::vector<sax::edge<std::string>> edges; //src, dest, weight(int)
std::cin >> edges;
std::string start, end;
std::cin >> start >> end;

std::unordered_set<std::string> all_nodes;
for(auto &edge : edges){
    graph[edge.src].push_back({edge.dest, edge.weight});
    all_nodes.insert(edge.src);
    all_nodes.insert(edge.dest);
}
for(auto &node : all_nodes){
    if(state[node] == 0){
        dfs(node);
    }
}
if(has_cycle){
    std::cout << "CYCLE\n";
    return 0;
}
std::reverse(topo.begin(), topo.end()); //=> A->B->C

const int INF = 1e9; //infinity
std::unordered_map<std::string, int> dist; //initially dont know the distance so set them to infinity
for(auto &node : all_nodes) dist[node] = INF;
dist[start] = 0; //base case

for(auto &u : topo){
    if(dist[u] != INF){
        for(auto &[v, w] : graph[u]){
            dist[v] = std::min(dist[v], dist[u] + w);
        }
    }
}
if(dist[end] == INF){
    std::cout << "NO PATH\n";
} else {
    std::cout << dist[end] << "\n";
}
return 0;
}

```

Time complexity: this algorithm has a time complexity of  $O(n + m)$ , because

the algorithm visits each node once in the DFS for topological sorting. Then, it relaxes each edge once in the shortest-path phase.

## Assignment 2 - Fire department placement

My code:

```
#include <iostream>
#include "utils.h" // for reading vectors
#include <vector>
#include <unordered_map>
#include <unordered_set>
#include <string>
#include <queue>
#include <algorithm>

int main() {
    /* TODO:
        A municipality wants to place a fire department in one of its neighborhoods such that the sum of distances from the fire department to all other neighborhoods is minimized.

        Write a program that computes its optimal placement and prints the smallest sum of distances from the fire department to all other neighborhoods.

        The program must read a list of edges representing an *undirected* graph from its standard input. Each edge is represented by a comma-separated list between square brackets, e.g. `[(A, B, 5), (B, C, 10), (C, D, 2)]` indicates that there is a road between neighborhoods A and B with length w.

        The program must then compute at which neighborhood the fire department should be placed such that the sum of distances from that neighborhood to all other neighborhoods is minimized.

        The time complexity of your solution must be `O(n m log n)`, where `n` is the number of nodes and `m` is the number of edges in the graph.
    */
    std::vector<sax::edge<std::string>> edges;
    std::cin >> edges;

    std::unordered_map<std::string, std::vector<std::pair<std::string, int>>> graph;
    std::unordered_set<std::string> all_nodes;

    for(auto& edge : edges){
        graph[edge.src].push_back({edge.dest, edge.weight});
        graph[edge.dest].push_back({edge.src, edge.weight});
        all_nodes.insert(edge.src);
        all_nodes.insert(edge.dest);
    }
}
```

```

}

const int INF = 1e9;
std::string best_node;
int best_sum = INF;

//dijkstra from each node
for(auto& start : all_nodes){
    std::unordered_map<std::string, int> dist;
    for(auto &node : all_nodes) dist[node] = INF;
    dist[start] = 0;

    //in Dijkstra's algorithm, we want the smallest distance (the "closest" node first)
    //So we need to flip the order to make it a min-heap - and that's what std::greater
    std::priority_queue<std::pair<int, std::string>, std::vector<std::pair<int, std::string>> pq;
    pq.push({0, start});

    while(!pq.empty()){
        auto [d, u] = pq.top(); //u: current node being visited //top() gives node wi
        pq.pop(); //d: current distance from start to u //pop() the sma
        if(d > dist[u]) continue; //we may input shorter path to the same node later.

        for(auto &[v, w] : graph[u]){
            if(dist[v] > dist[u] + w){
                dist[v] = dist[u] + w;
                pq.push({dist[v], v});
            }
        }
    }

    int total = 0;
    for(auto &[n, d] : dist){
        if(d != INF) total += d;
    }

    if(total < best_sum){
        best_sum = total;
        best_node = start;
    }
}

std::cout << best_node << ", " << best_sum << "\n";
return 0;
}

// Start      To A      To B      To C      Total
// A          0          2          5          7
// B          2          0          3          5

```

```

// C      5      3      0      8

// Start = A => total = 7
// => best_sum = 7, best_node = "A"

// Start = B => total = 5
// => 5 < 7 => update: best_sum = 5, best_node = "B"

// Start = C => total = 8
// => 8 > 5 => no update

```

Time complexity: this algorithm has a time complexity of  $O(n * m \log n)$ , because the program runs Dijkstra's algorithm once for every node in the graph. One Dijkstra run takes  $O(m \log n)$  time using a priority queue (`std::priority_queue`). Since we repeat it for each node ( $n$  times)  $\Rightarrow$  total =  $O(n * m \log n)$ .

### Assignment 3 - Undirected cycle detection

My code:

```

#include <iostream>
#include <unordered_map>
#include "utils.h"
#include <vector>
#include <string>

//DSU structure
std::unordered_map<std::string, std::string> parent;
std::unordered_map<std::string, int> size;

//find function with path compression
std::string find(const std::string &x){
    if(parent[x] == x) return x;
    return parent[x] = find(parent[x]); //path compression
}

bool union_set(const std::string &a, const std::string &b){
    std::string rootA = find(a);
    std::string rootB = find(b);
    if(rootA == rootB) return false; //cuz cycle detected //if both node have the same root
                                    //so adding edge between them would make cycle
    if(size[rootA] < size[rootB]){
        std::swap(rootA, rootB);
    }
    parent[rootB] = rootA;
    size[rootA] += size[rootB];
}

```

```

        return true;
    }

int main() {
    /* TODO:
       Write a program that reads a list of edges representing an *undirected* graph from
       input (given as a comma-separated list between square brackets, e.g. `[(A, B), (B, C)]`).
       The program must add these edges, one by one, to an initially empty graph, and before
       each edge, it must determine whether that edge would introduce a cycle into the graph.
       In case the edge would introduce a cycle, the edge must not be added to the graph and
       the program must stop processing further edges. Otherwise, the edge is added to the graph and
       the program continues with the next edge.
       The program must then print how many edges were added before the first cycle was created.
       Your program must run in `O(m log(n))` time, where `n` is the number of nodes and `m` is
       the number of edges.
    */
    //didn't use sax::edge cuz we dont need weights
    std::vector<std::pair<std::string, std::string>> edges;
    std::cin >> edges;
    int added = 0;

    for(auto &edge : edges){
        //initialize parent and size if new node
        if(!parent.count(edge.first)){
            parent[edge.first] = edge.first;
            size[edge.first] = 1;
        }
        if(!parent.count(edge.second)){
            parent[edge.second] = edge.second;
            size[edge.second] = 1;
        }
        if(!union_set(edge.first, edge.second)){
            //cycle detected
            break;
        }
        added++;
    }
    std::cout << added << "\n";
    return 0;
}

```

Time complexity: this algorithm has a time complexity of  $O(m \log n)$ , because each edge is processed once  $\Rightarrow m$  times total. For each edge, the algorithm performs `find()` and `union()` operations on the disjoint set. With path compression and union by size, each operation runs in almost constant time  $\Rightarrow$  technically  $O(\log n)$  (or even tighter,  $O(\alpha(n))$ ), where  $\alpha$  is the inverse Ackermann function).

## Assignment 4 - Minimum satellite communication network

My code:

```
#include <iostream>
#include "utils.h"
#include "coordinate.h"
#include <unordered_map>
#include <cmath>
#include <algorithm>
#include <vector>
#include <edge.h>

// Euclidean Distance
// For two coordinates (x1, y1, z1) and (x2, y2, z2):
// sqrt((x1 - x2)^2 + (y1 - y2)^2 + (z1 - z2)^2)

// Kruskal's algorithm steps:
// Compute all possible edges (for every pair of satellites).
// Sort edges by weight (distance).
// Initialize DSU (each satellite is its own parent).
// For each edge (smallest first):
// If connecting them doesn't form a cycle (=>union_set() returns true), add its weight
// Stop when you have added n - 1 edges.

//cannot use sax::edge cuz it's always int weight
struct Edge{
    int src, dest;
    double weight;
};

// We use std::vector<int> parent, size; instead of unordered_map because:
// Satellite nodes are numbered 0 to n - 1, not arbitrary names.
// Vectors allow direct fast access ( $O(1)$ ), while maps need hash lookups (slower).
// We can initialize all nodes easily with a simple loop.
// std::unordered_map<int, int> parent;
```

```

// std::unordered_map<int, int> size;
std::vector<int> parent, size;

int find(int x){
    if(parent[x] == x) return x;
    return parent[x] = find(parent[x]); //path compression
}

bool union_set(int a, int b){
    int rootA = find(a);
    int rootB = find(b);
    if(rootA == rootB) return false; //cycle detected

    if(size[rootA] < size[rootB]){
        std::swap(rootA, rootB);
    }
    parent[rootB] = rootA;
    size[rootA] += size[rootB];
    return true;
}

// //compute Euclidean distance between two coordinates
// double distance(const sax::coordinate &a, const sax::coordinate &b){
//     double dx = a.x - b.x;
//     double dy = a.y - b.y;
//     double dz = a.z - b.z;
//     return std::sqrt(dx*dx + dy*dy + dz*dz);
// }

int main()
{
    /* TODO:
       A communication network consists of a number of satellites that need to be connected.
       Satellites can be connected directly to each other, or indirectly through other satellites.
       how the internet works.

       Since communication is costly, the network must be designed such that the total cost of connecting all the satellites is minimized. Also, the network must not contain any cycles, since that would waste resources. In other words, there should be just one unique communication path between any two satellites.

       Given (x, y, z) coordinates of satellites, write a program that computes the minimum cost of connecting all the satellites in a communication network, where the cost to connect two satellites is given by the Euclidean distance between them.

       The program receives a list of satellites from its standard input (given as a comma-separated list of coordinates enclosed in square brackets, e.g. `[(0,0,0),(1,1,1),(2,2,2)]`), where each satellite is represented by its (x, y, z) coordinates.
    */
}

```

coordinates. Use the `sax::coordinate` type to represent the coordinates of a satellite.

The program must then compute the minimum cost to connect all satellites in a community. Print this cost to its standard output with two digits after the decimal point (e.g. `3.42`).

```
/*
std::vector<sax::coordinate> satellites;
std::cin >> satellites;

int n = satellites.size();
if(n <= 1){
    std::cout << "0.00\n";
    return 0;
}

std::vector<Edge> edges;

for(int i = 0; i < n; i++){
    for(int j = i + 1; j < n; j++){
        edges.push_back({i, j, satellites[i].distance_to(satellites[j])});
    }
}
//ex: satellites[0] = (0, 0, 0) => index = 0
//    satellites[1] = (1, 0, 0) => index = 1
//    satellites[2] = (0, 1, 0) => index = 2

//edges.push_back({0, 1, satellites[0].distance_to(satellites[1])});
//=> Edge created: {src=0, dest=1, weight=1.00}

//edges.push_back({0, 2, satellites[0].distance_to(satellites[2])});
//=> Edge created: {src=0, dest=2, weight=1.00}

//edges.push_back({1, 2, satellites[1].distance_to(satellites[2])});
//=> Edge created: {src=1, dest=2, weight=1.41}

//sort edges by weight
std::sort(edges.begin(), edges.end(), [] (const Edge &a, const Edge &b){
    return a.weight < b.weight;
});

//using vector not unordered_map so need to change to this
parent.resize(n);           //allocate space for all n satellites
size.resize(n, 1);           //size is initialized to 1 for each node
//DSU - set parent and size for later use in union_set and find
for(int i = 0; i < n; i++){
    parent[i] = i;
    // size[i] = 1;
```

```

    }

    double total = 0.0;
    int edges_used = 0;

    // Kruskal's Algorithm
    // Kruskal doesn't care about the coordinates, only about the indices (0,1,2).
    // It uses these indices in the DSU (Union-Find) to merge sets:
    // Connect 0-1
    // Connect 0-2
    // Don't connect 1-2 (would make a cycle )
    // Minimum total = 1.00 + 1.00 = 2.00

    for(auto &edge : edges){
        if(union_set(edge.src, edge.dest)){
            total += edge.weight;
            edges_used++;

            if(edges_used == n-1){
                break; //MST complete
            }
        }
    }
    std::cout.setf(std::ios::fixed);
    std::cout.precision(2);
    std::cout << total << "\n";
    return 0;
}

```

I KEEP INSPECTING ERROR IN THIS TEST (I've also asked some of my friend they got the same problem): The answer for the following test input (line 3) is incorrect: [(0,0,0),(1,2,3),(8,5,2),(2,3,0),(5,5,5)] got 16.55 expected 22.71

Time complexity: this algorithm has a time complexity of ....., because .....

## Assignment 5 - Shortest paths on minimum spanning trees

My code:

```

#include <iostream>
#include "utils.h"
#include "coordinate.h"
#include <unordered_map>
#include <cmath>
#include <algorithm>

```

```

#include <vector>
#include <edge.h>
#include <string>

std::unordered_map<std::string, std::string> parent;
std::unordered_map<std::string, int> size;

std::string find(const std::string &x){
    if(parent[x] == x) return x;
    return parent[x] = find(parent[x]);
}

bool union_set(const std::string &a, const std::string &b){
    std::string rootA = find(a);
    std::string rootB = find(b);
    if(rootA == rootB) return false;

    if(size[rootA] < size[rootB]){
        std::swap(rootA, rootB);
    }
    parent[rootB] = rootA;
    size[rootA] += size[rootB];
    return true;
}

// In an MST or any undirected graph, there can be multiple routes between two nodes.
// The old DFS stopped after the first one it found,
// but the new DFS tries all paths and keeps the smallest cost.
void dfs(const std::string &current, const std::string &target, std::unordered_map<std::string, std::unordered_map<std::string, bool>> &graph, std::unordered_map<std::string, bool> &visited, int current_cost, int &min_cost){
    if(current == target){
        min_cost = std::min(min_cost, current_cost);
        return;
    }
    visited[current] = true;
    for(auto &[neighbor, w] : graph[current]){
        if(!visited[neighbor]){
            dfs(neighbor, target, graph, visited, current_cost + w, min_cost);
        }
    }
    visited[current] = false;
}

int main() {
    /* TODO:
       Given an undirected graph, write a program that computes the minimum spanning tree */
}

```

either Prim's or Kruskal's algorithm, and then computes the shortest path (i.e., with minimum cost) between two given nodes on the minimum spanning tree.

The program receives a list of edges representing an \*undirected weighted\* graph from the input (given as a comma-separated list between square brackets, e.g. `[(A, B, 5), (B, C, 10), (C, D, 2)]` followed by two characters representing the start and end nodes (e.g. `A D`) – similar to the problem of this week.

The program's output must be the cost of the shortest path between the start and end nodes, or `NO PATH` if there is no path between the start and end nodes.

```
/*
//ex: [(A,B,5),(B,C,10),(C,D,2)] A D
std::vector<sax::edge<std::string>> edges;
std::cin >> edges;
std::string start, end;
std::cin >> start >> end;

for(auto &edge : edges){
    if(!parent.count(edge.src)){
        parent[edge.src] = edge.src; //set parent and size for later use in union_set
        size[edge.src] = 1;
    }
    if(!parent.count(edge.dest)){
        parent[edge.dest] = edge.dest;
        size[edge.dest] = 1;
    }
}
//Kruskal's algorithm -> build MST
std::sort(edges.begin(), edges.end(), [](auto &a, auto &b){
    if (a.weight != b.weight) return a.weight < b.weight;
    if (a.src != b.src) return a.src < b.src;
    return a.dest < b.dest;
});

std::unordered_map<std::string, std::vector<std::pair<std::string, int>>> graph;
for(auto &edge : edges){
    if(union_set(edge.src, edge.dest)){ //build graph after union_set to avoid cycle
        graph[edge.src].push_back({edge.dest, edge.weight});
        graph[edge.dest].push_back({edge.src, edge.weight});
    }
}
std::unordered_map<std::string, bool> visited;
const int INF = 1e9;
int min_cost = INF;

dfs(start, end, graph, visited, 0, min_cost);
```

```

if (min_cost != INF){
    std::cout << min_cost << std::endl;
}
else{
    std::cout << "NO PATH" << std::endl;
}
return 0;
}

```

Time complexity: this algorithm has a time complexity of  $O(m \log n)$ , because when building the MST (Kruskal's algorithm): Sorting all  $m$  edges takes  $O(m \log m)$ , which is roughly  $O(m \log n)$  since  $m \leq n^2$ . Each union\_set and find operation takes almost constant time ( $O(\log n)$ ) due to path compression and union by size. So total for MST =  $O(m \log n)$ . while the DFS traversal of the MST only adds linear time => the MST has exactly  $n - 1$  edges => so DFS runs in  $O(n + (n - 1)) = O(n)$  time.