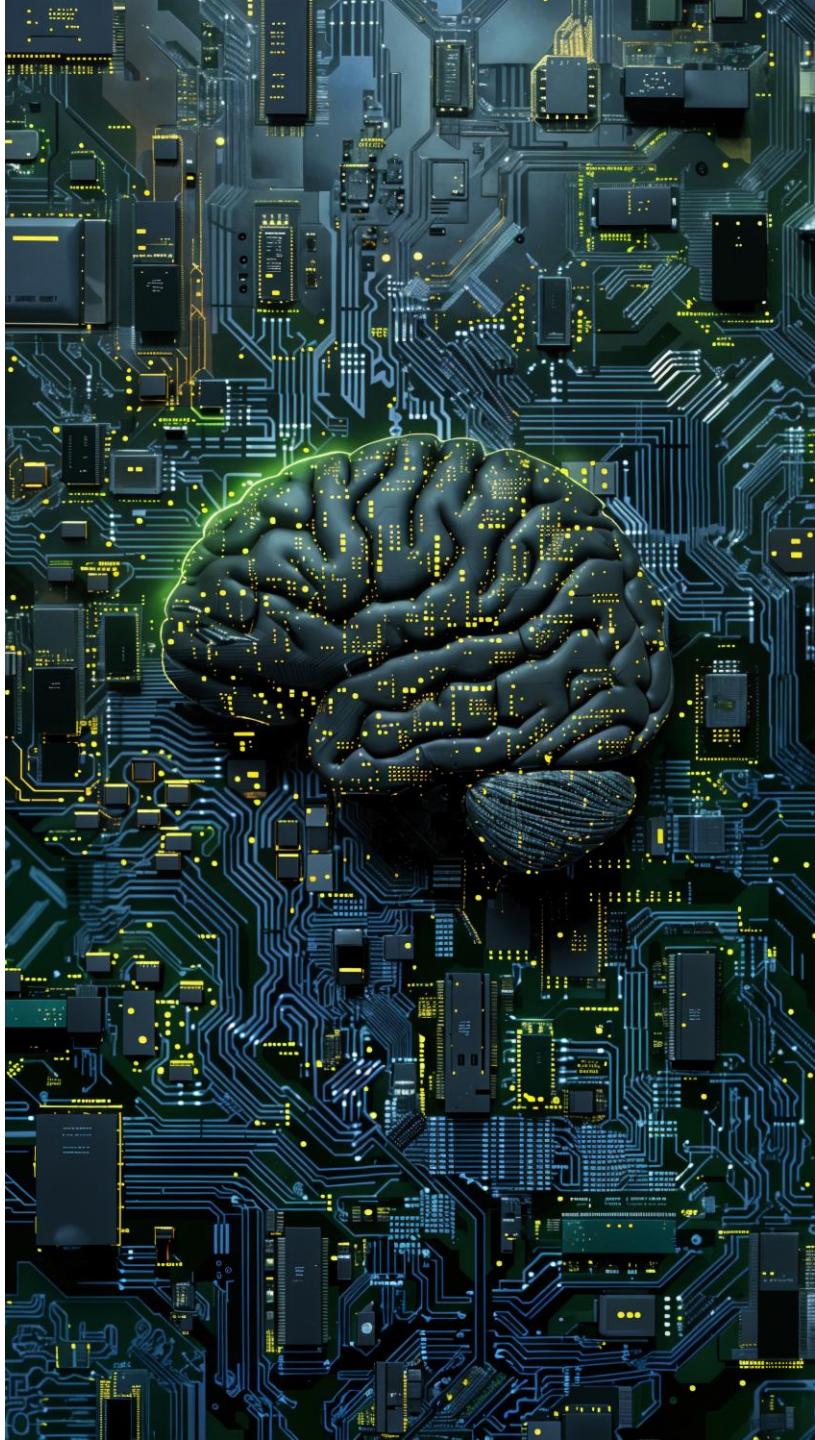


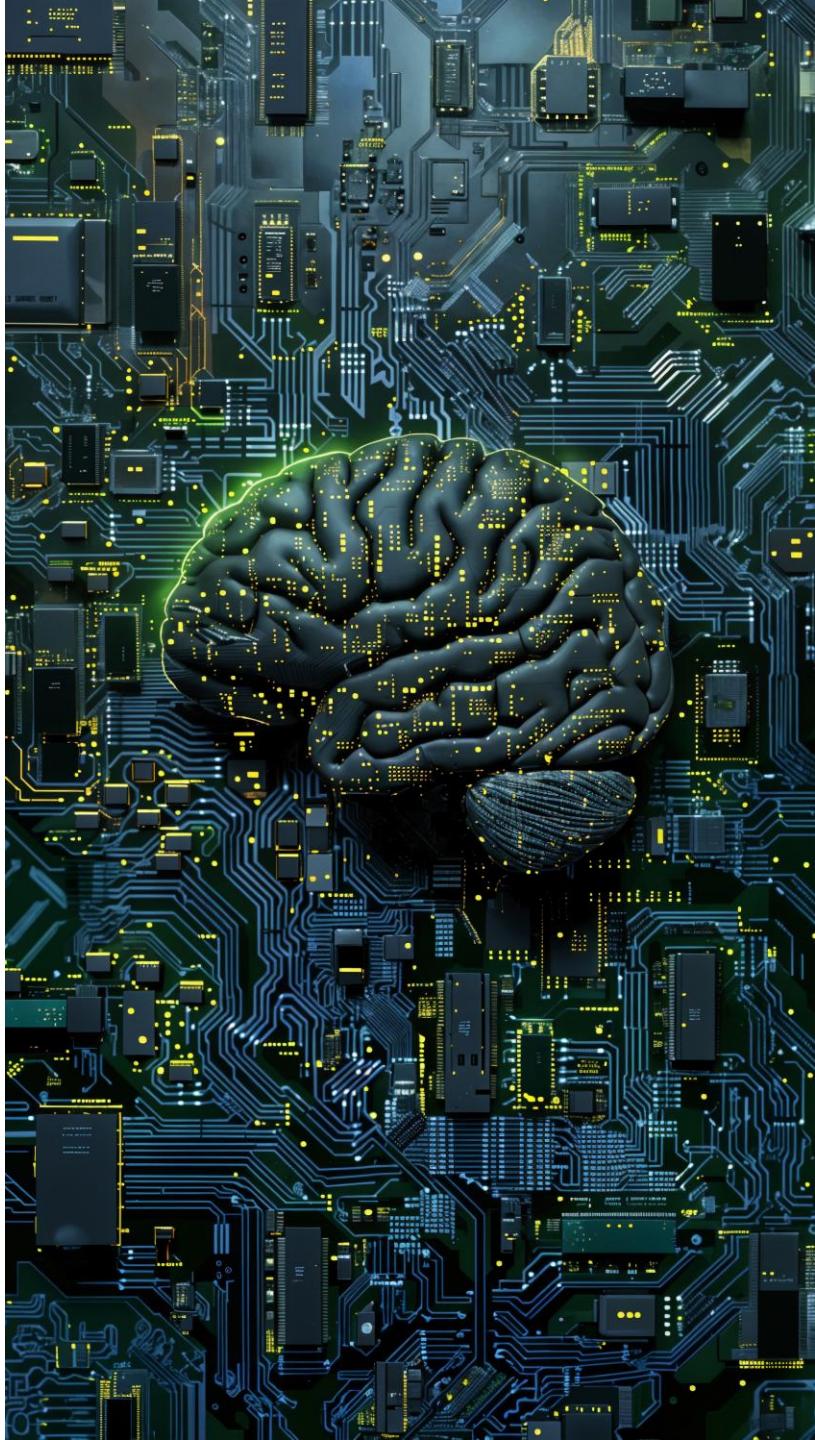
# Machine Learning

Week 3 – Working with data, classification



# Working with data

Types of data, augmenting, filling, transforming



# Pandas is your friend

Most of the data formats can be read by pandas:

```
import pandas as pd  
  
df = pd.read_csv("dataset.csv", sep=",")
```

Pandas works with **DataFrame**'s

# Pandas is your friend

Looking at the data

```
data.head() #take a look at the data
```

...	HSP	EA	# EB	# EC	# SP
0	N	Below expectations	6.2	Missing value	21.0
1	E	Below expectations	Missing value	Missing value	Missing value
2	N	Below expectations	14.6	13.0	33.0
3	K	Missing value	Missing value	48.3	13.0
4	N	Below expectations	36.6	Missing value	3.0

A green box highlights the '# EB' column, labeled 'feature'. A red box highlights the second row (index 1), labeled 'observation'.

# Pandas is your friend

Looking at the data

```
data.head() #take a look at the data
```

...	HSP	EA	# EB	# EC	# SP
0	N	Below expectations	6.2	Missing value	21.0
1	E	Below expectations	Missing value	Missing value	Missing value
2	N	Below expectations	14.6	13.0	33.0
3	K	Missing value	Missing value	48.3	13.0
4	N	Below expectations	36.6	Missing value	3.0



index    categorical features    numerical features

# Pandas is your friend

Selecting rows (observations) by **Row Index**

`df.loc[4]`

↳ 4	
HSP	N
EA	Below expectations
EB	36.6
EC	<i>Missing value</i>
SP	3.0

# Pandas is your friend

## Selecting columns

```
df.loc[:, "EA"]  
df[ "EA" ]
```

	EA	...
0	Below expectations	
1	Below expectations	
2	Below expectations	
3	<i>Missing value</i>	
4	Below expectations	
5	Meets expectations	
6	<i>Missing value</i>	
7	Sufficient	
8	Below expectations	
9	Good	

78 rows x 1 cols 10 ▾ per page

```
df.loc[:, [ "SP", "HSP" ]]  
df[ [ "SP", "HSP" ] ]
```

	...	# SP	HSP
0		21.0	N
1	<i>Missing value</i>	E	
2		33.0	N
3		13.0	K
4		3.0	N
5		24.0	E
6		45.0	G
7		28.0	E
8		46.0	K
9		11.0	K

78 rows x 2 cols 10 ▾ per page

# Pandas is your friend

Combining row and column selection

```
df.loc[5:9, "EB":]
```

...	# EB	# EC	# SP
5	0.8	<i>Missing value</i>	24.0
6	62.4	<i>Missing value</i>	45.0
7	24.0	4.0	28.0
8	4.6	51.0	46.0
9	<i>Missing value</i>	7.0	11.0

# Pandas is your friend

Row and column selection by 0-based (array-like) indexing

```
df.iloc[5:9, 2:]
```

...	# EB	# EC	# SP
5	0.8	<i>Missing value</i>	24.0
6	62.4	<i>Missing value</i>	45.0
7	24.0	4.0	28.0
8	4.6	51.0	46.0

# Pandas is your friend

Selecting data based on condition

```
# select items where SP <= 40 and HSP taking values of N or E  
df[(df["SP"] <= 40) & df['HSP'].isin(['N', 'E']) ]
```

...	HSP	EA	# EB	# EC	# SP
0	N	Below expectations	6.2	<i>Missing value</i>	21.0
2	N	Below expectations	14.6	13.0	33.0
4	N	Below expectations	36.6	<i>Missing value</i>	3.0
5	E	Meets expectations	0.8	<i>Missing value</i>	24.0
7	E	Sufficient	24.0	4.0	28.0

# Pandas: basic cleaning

Looking at the data once more

```
data.head() #take a look at the data
```

...	HSP	EA	# EB	# EC	# SP
0	N	Below expectations	6.2	Missing value	21.0
1	E	Below expectations	Missing value	Missing value	Missing value
2	N	Below expectations	14.6	13.0	33.0
3	K	Missing value	Missing value	48.3	13.0
4	N	Below expectations	36.6	Missing value	3.0

Missing value  
**NaN** (for numerical)  
Null (for others)

# Pandas: basic cleaning

How many missing values do we have?

```
print(data.isna().sum())
```

	...	# 0
HSP		0
EA		3
EB		4
EC		9
SP		1

```
df = df.dropna()  
# or in-place  
df.dropna(inplace=True)
```

} Just remove observations  
that have missing features

# Pandas: basic cleaning

Based on domain knowledge (if you know what your data is):

- EA, EB, EC are exam grades
- SP are study points

replace missing values with 0's:

```
# replace NaNs in EB, EC, SP with 0
df[["EB", "EC", "SP"]] = df[["EB", "EC", "SP"]].fillna(0)
```

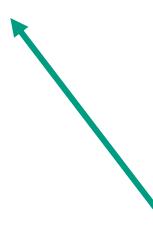
```
# replace NaNs in EA with "Below expectations"
df["EA"] = df["EA"].fillna("Below expectations")
```

```
# Or perform in-place replacement
df.fillna({"EA": "Below expectations"}, inplace=True)
```

# It's better to use scikit-learn for this!

```
from sklearn.impute import SimpleImputer  
  
imputer = SimpleImputer(strategy="constant", fill_value=0)  
  
df[:] = imputer.fit_transform(df[:])
```

	HSP		EA	EB	EC	SP
0	N	Below expectations	6.2	0.0	21.0	
1	E	Below expectations	0.0	0.0	0.0	
2	N	Below expectations	14.6	13.0	33.0	
3	K		0	0.0	48.3	13.0
4	N	Below expectations	36.6	0.0	3.0	

- 
- strategies:
- constant
  - mean
  - median
  - most\_frequent

## It's better to use scikit-learn for this!

```
pipeline = Pipeline(steps=[  
    ("imputer", ColumnTransformer(  
        transformers=[  
            ("numeric", SimpleImputer(strategy="constant", fill_value=0), ["EB", "EC", "SP"]),  
            ("other", SimpleImputer(strategy="constant", fill_value="Below expectations"), ["EA"]),  
        ],  
        remainder="passthrough"  
    )),  
    ("scaler", MinMaxScaler()),  
    ("regressor", LinearRegression()),  
])
```

# Pandas: knowing data

	<b>HSP</b>		<b>EA</b>	<b>EB</b>	<b>EC</b>	<b>SP</b>
0	N	Below expectations	6.2	0.0	21.0	
1	E	Below expectations	0.0	0.0	0.0	
2	N	Below expectations	14.6	13.0	33.0	
3	K	Below expectations	0.0	48.3	13.0	
4	N	Below expectations	36.6	0.0	3.0	
5	E	Meets expectations	0.8	0.0	24.0	
6	G	Below expectations	62.4	0.0	45.0	
7	E	Sufficient	24.0	4.0	28.0	
8	K	Below expectations	4.6	51.0	46.0	
9	K	Good	0.0	7.0	11.0	

# Pandas: knowing data

Transforming categorical data:

- Ordered data: `OrdinalEncoder`
- Unordered data: `OneHotEncoder`
- Target variable: `LabelEncoder`

# Pandas: knowing data

	HSP		EA	EB	EC	SP
0	N	Below expectations	6.2	0.0	21.0	
1	E	Below expectations	0.0	0.0	0.0	
2	N	Below expectations	14.6	13.0	33.0	
3	K	Below expectations	0.0	48.3	13.0	
4	N	Below expectations	36.6	0.0	3.0	
5	E	Meets expectations	0.8	0.0	24.0	
6	G	Below expectations	62.4	0.0	45.0	
7	E	Sufficient	24.0	4.0	28.0	
8	K	Below expectations	4.6	51.0	46.0	
9	K	Good	0.0	7.0	11.0	

EA is ordered

1. "Below expectations"
2. "Sufficient"
3. "Meets expectations"
4. "Good"
5. "Exceeds expectations"

HSP seems unordered

'N', 'E', 'K', 'G', 'B'

## Encoding ordered categorical data

```
from sklearn.preprocessing import OrdinalEncoder  
  
ordinal_encoder = OrdinalEncoder(  
    categories=[["Below expectations",  
                "Sufficient",  
                "Meets expectations",  
                "Good",  
                "Excellent"]])  
  
df["EA_"] = ordinal_encoder.fit_transform(df[["EA"]])
```

	EA	EA_
0	Below expectations	0.0
1	Below expectations	0.0
2	Below expectations	0.0
3	Below expectations	0.0
4	Below expectations	0.0
5	Meets expectations	2.0
6	Below expectations	0.0
7	Sufficient	1.0
8	Below expectations	0.0
9	Good	3.0

## Encoding ordered categorical data (pipeline variant)

```
pipeline = Pipeline(steps=[  
    ("imputer", ColumnTransformer( ...)),  
    ("ordinal_encoder", ColumnTransformer(  
        transformers=[  
            ("EA", OrdinalEncoder(categories=[["Below expectations",  
                                              "Sufficient",  
                                              "Meets expectations",  
                                              "Good",  
                                              "Excellent"]]), ["EA"])),  
        ],  
    remainder="passthrough"  
),  
...  
])
```

## Encoding unordered categorical data

```
from sklearn.preprocessing import OneHotEncoder

onehot_encoder = OneHotEncoder(sparse_output=False, drop="if_binary")

hsp_encoded = onehot_encoder.fit_transform(df[["HSP"]])

hsp_encoded_df = pd.DataFrame(hsp_encoded,
                               columns=onehot_encoder.get_feature_names_out(["HSP"]))

pd.concat([df["HSP"], hsp_encoded_df], axis=1).head()
```

## Encoding unordered categorical data

```
from sklearn.preproce      HSP   HSP_B   HSP_E   HSP_G   HSP_K   HSP_N  
onehot_encoder = OneHotEncoder(sparse=False)  
hsp_encoded = onehot_encoder.fit_transform(df[["HSP"]])  
hsp_encoded_df = pd.DataFrame(hsp_encoded, columns=["HSP_B", "HSP_E", "HSP_G", "HSP_K", "HSP_N"], index=df.index)  
pd.concat([df[["HSP"]], hsp_encoded_df], axis=1)
```

	HSP	HSP_B	HSP_E	HSP_G	HSP_K	HSP_N
0	N	0.0	0.0	0.0	0.0	1.0
1	E	0.0	1.0	0.0	0.0	0.0
2	N	0.0	0.0	0.0	0.0	1.0
3	K	0.0	0.0	0.0	1.0	0.0
4	N	0.0	0.0	0.0	0.0	1.0
5	E	0.0	1.0	0.0	0.0	0.0
6	G	0.0	0.0	1.0	0.0	0.0
7	E	0.0	1.0	0.0	0.0	0.0
8	K	0.0	0.0	0.0	1.0	0.0
9	K	0.0	0.0	0.0	1.0	0.0

## Encoding unordered categorical data (pipeline variant)

```
pipeline = Pipeline(steps=[  
    ("imputer", ColumnTransformer( ...)),  
    ("ordinal_encoder", ColumnTransformer( ...)),  
  
    ("onehot_encoder", ColumnTransformer(  
        transformers=[  
            ("HSP", OneHotEncoder(sparse_output=True, drop="if_binary"), ["HSP"]),  
        ],  
        remainder="passthrough"  
    )),  
    ...  
])
```

## Take a deeper look at the data: distributions

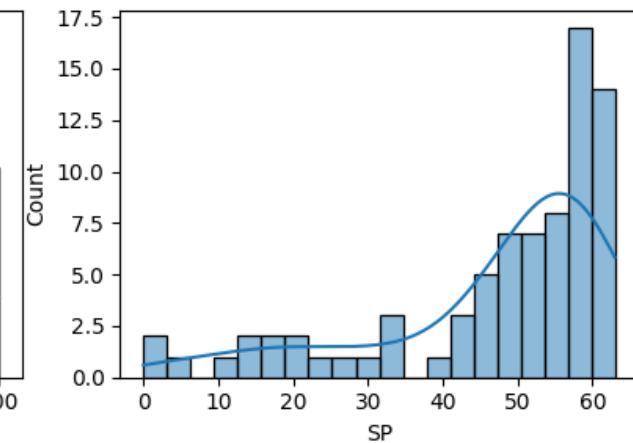
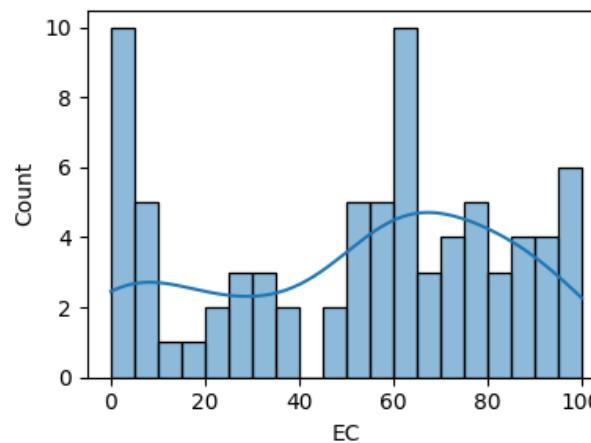
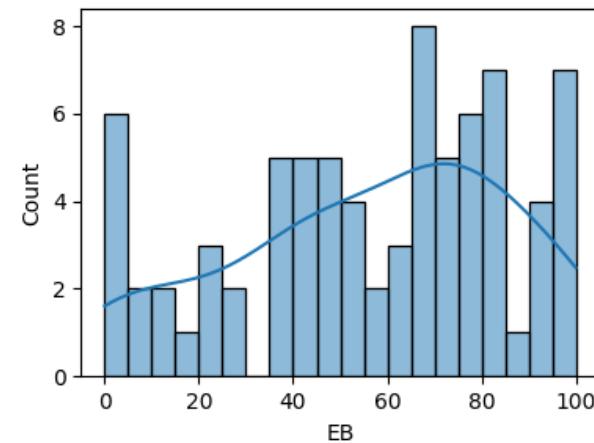
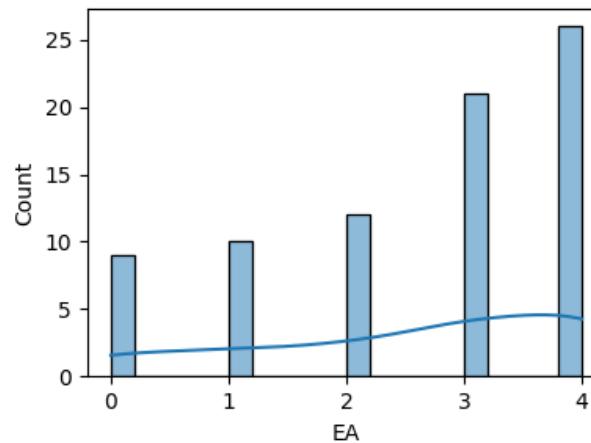
```
import seaborn as sns  
import matplotlib.pyplot as plt
```



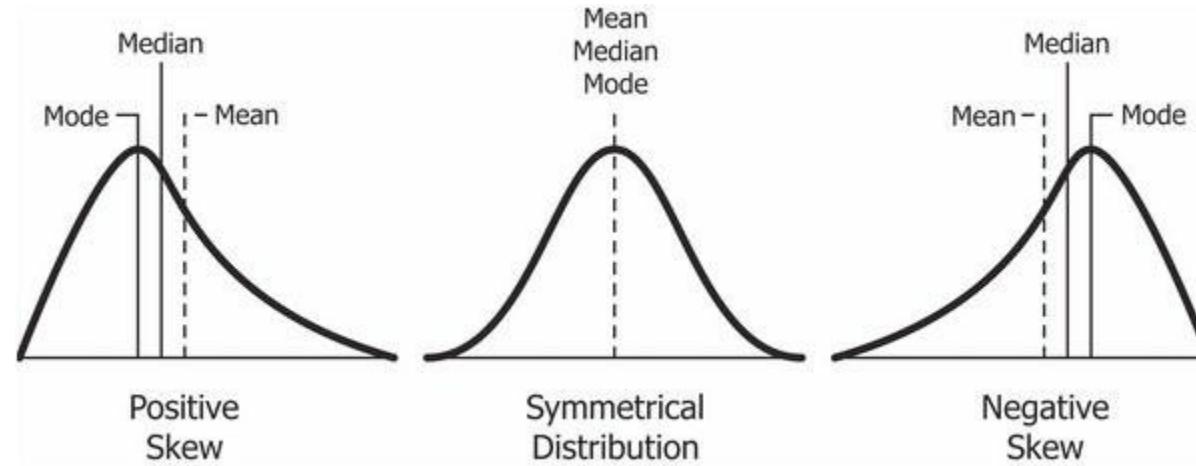
```
# plot histograms only for EA, EB, EC, SP  
fig, axs = plt.subplots(ncols=2, nrows=2, figsize=(8, 6))  
index = 0  
axs = axs.flatten()  
for k, v in df[['EA', 'EB', 'EC', 'SP']].items():  
    sns.histplot(v, ax=axs[index], bins=20, bins=True)  
    index += 1  
plt.tight_layout()
```



# Take a deeper look at the data: distributions



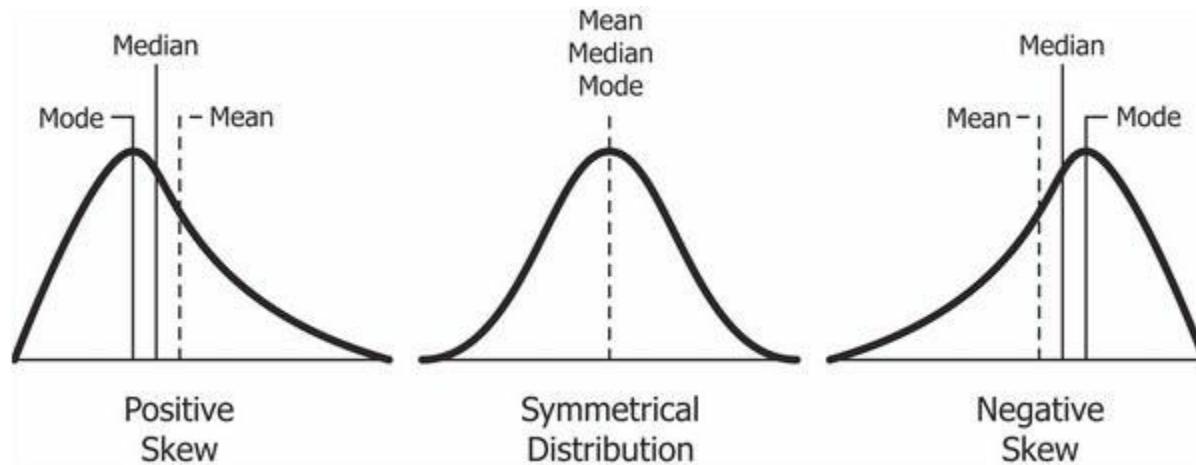
# Skewness



`df.skew()`

EA	-0.61
EB	-0.40
EC	-0.36
SP	-1.46

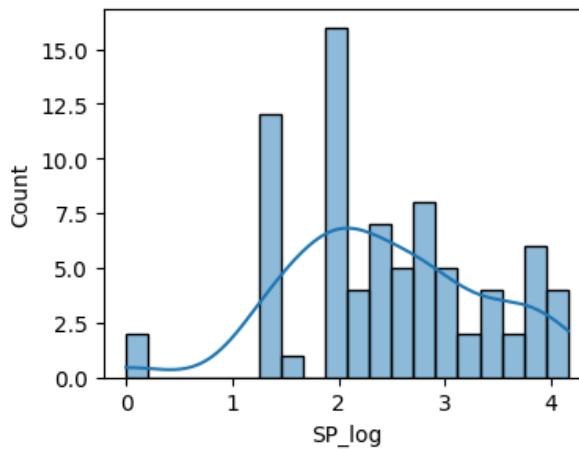
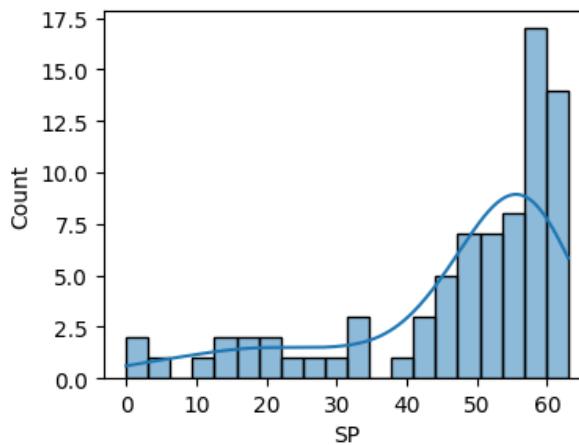
## Correcting skewness



```
# correct the negative skewness of NEG_SKEW using log transformation  
df['NEG_SKEW_log'] = np.log1p(df['NEG_SKEW'].max() - df['NEG_SKEW'])
```

```
# correct the positive skewness of POS_SKEW using log transformation  
df['POS_SKEW_log'] = np.log1p(df['POS_SKEW'])
```

# Correcting

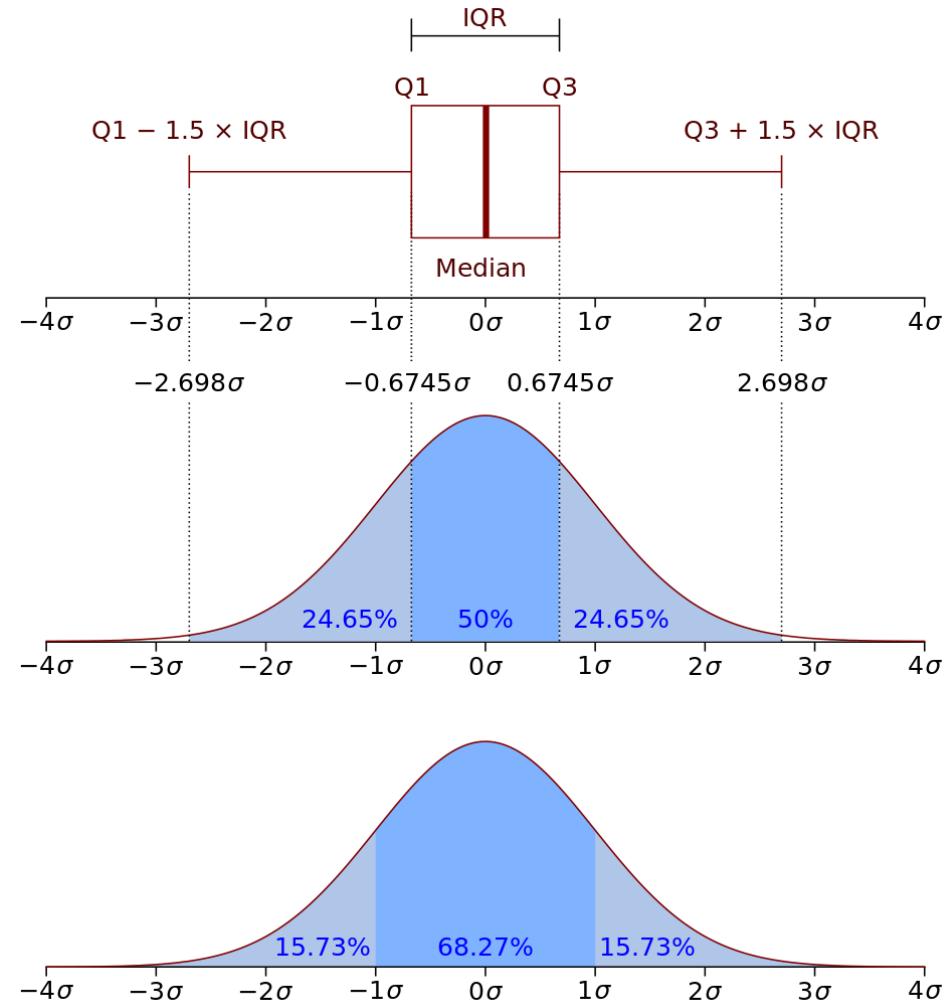


# It's better to use scikit-learn for this!

```
pipeline = Pipeline(steps=[  
    ("imputer", ColumnTransformer(...)),  
  
    ("scaler", MinMaxScaler()), # scale the data to 0..1, so we don't have problems with log's  
  
    ("log_transformer", ColumnTransformer( transformers=[  
  
        ("COL_NEG", FunctionTransformer(func=lambda x: np.log1p(1.0 - x)), ["COL_NEG"]),  
  
        ("COL_POS", FunctionTransformer(func=lambda x: np.log1p(x)), ["COL_POS"]),  
    ],  
    remainder="passthrough"  
)),  
    ("regressor", LinearRegression()),  
])
```

# Boxplots

← Interquartile range



# Boxplots

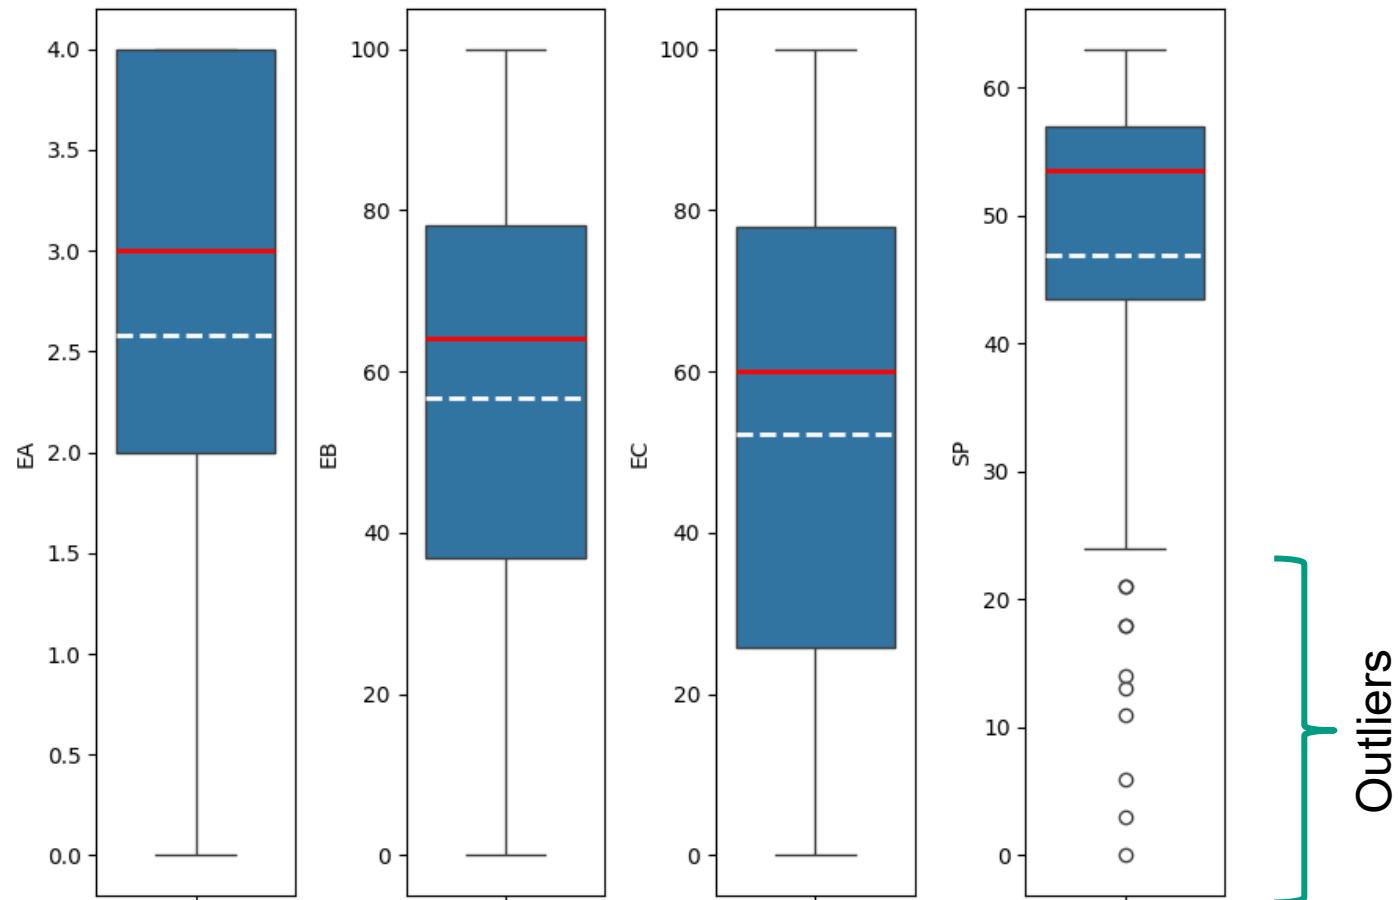
```
fig,axs = plt.subplots(ncols=4, nrows=1, figsize=(8, 6))
index = 0
axs = axs.flatten()
for k, v in df[['EA', 'EB', 'EC', 'SP']].items():

    sns.boxplot(y=k, data=df, ax=axs[index],
                 medianprops={"color": "red", "linewidth": 2},
                 meanprops={"color": "white", "linewidth": 2},
                 meanline=True, showmeans=True)

    index += 1
plt.tight_layout()
```



## Boxplots: outliers



# Violin plots

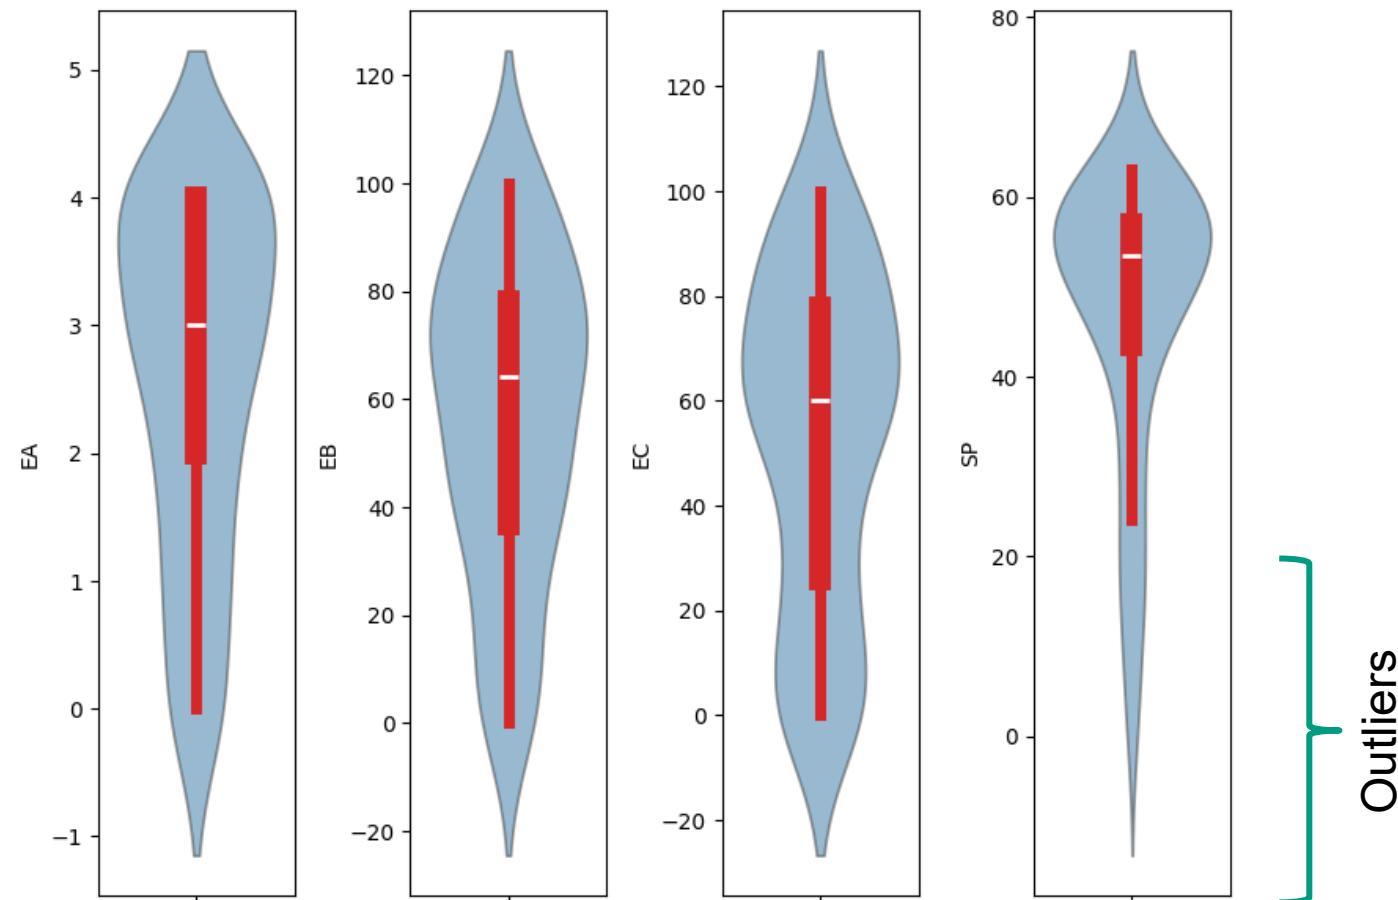
```
fig, axs = plt.subplots(ncols=4, nrows=1, figsize=(8, 6))
index = 0
axs = axs.flatten()
for k, v in df[['EA', 'EB', 'EC', 'SP']].items():

    sns.violinplot(y=k, data=df, ax=axs[index], alpha=0.5,
                    inner_kws={
                        "box_width": 10,
                        "whis_width": 5,
                        "color": sns.color_palette()[3]})

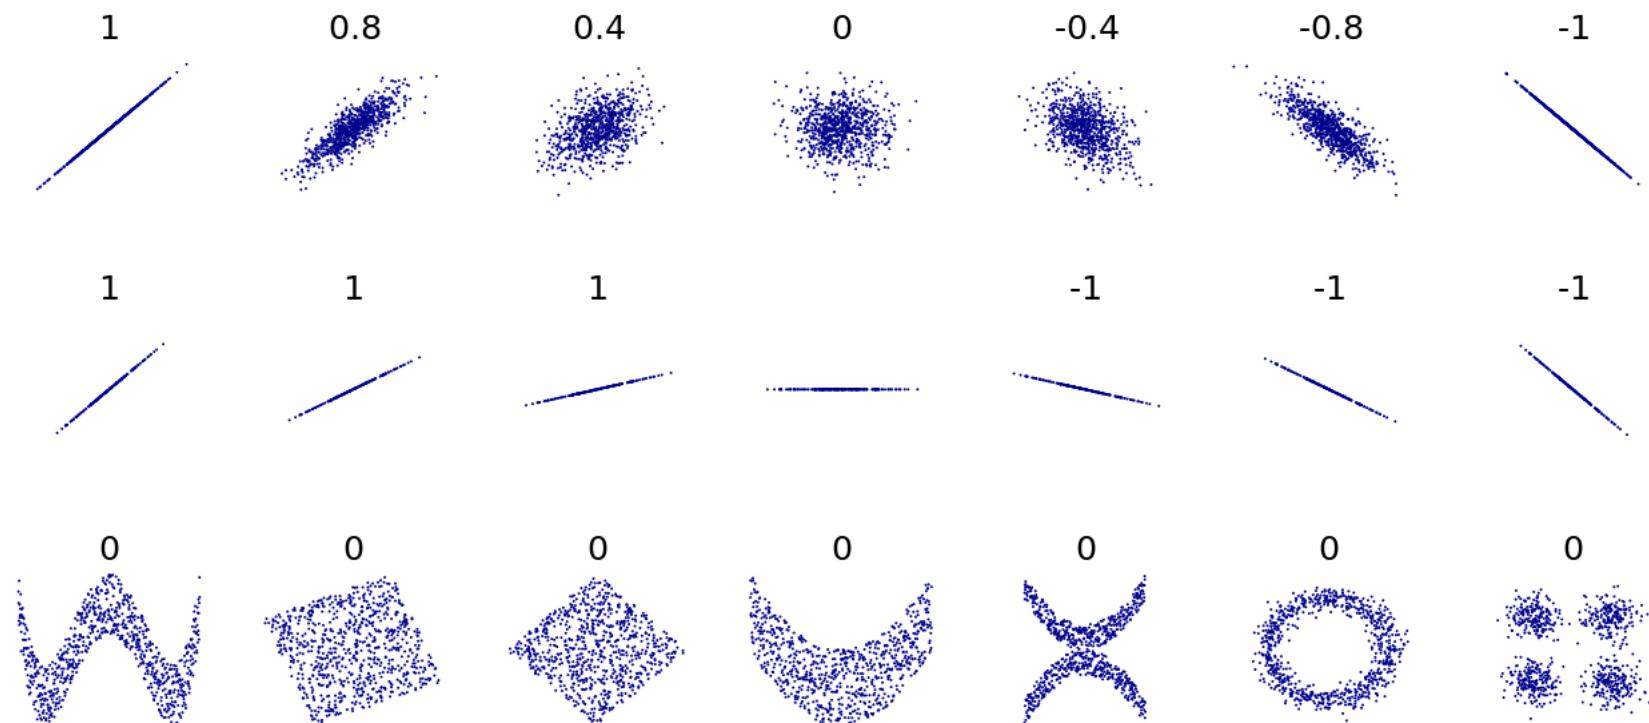
    index += 1
plt.tight_layout()
```



# Violin plots: outliers

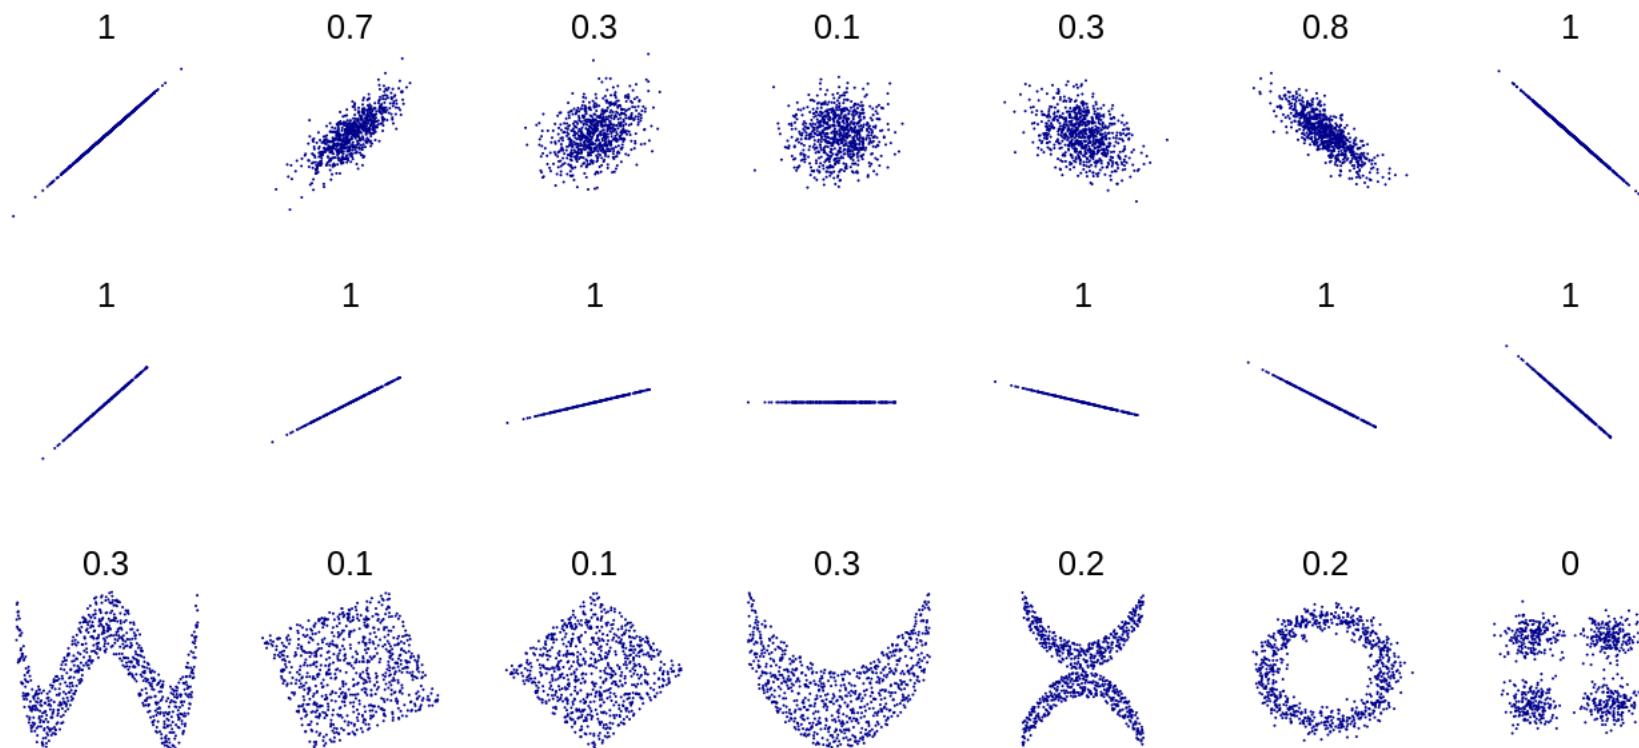


## Pearson correlation coefficient



Works only for linear ( $y=ax+b$ ) relationships!

## Other correlations: distance correlation



To get it: **pip install dcor**  
<https://dcor.readthedocs.io/en/stable/index.html>

# Looking at correlations

```
pd.set_option('display.precision', 2)

pd.DataFrame({
    'Pearson': df.corr(method='pearson')['SP'],
    'Spearman': df.corr(method='spearman')['SP']
})
```

## Correlation:

- ~0: none at all
- <0.5: none ÷ weak
- >0.5: it seems data is correlated
- >0.8: strong

# Looking at correlations

	Pearson	Spearman
EA	0.55	0.64
EB	0.72	0.71
EC	0.80	0.81
SP	1.00	1.00
HSP_B	0.62	0.75
HSP_E	-0.38	-0.28
HSP_G	0.07	-0.17
HSP_K	-0.52	-0.51
HSP_N	-0.35	-0.27

## Correlation:

~0: none at all

<0.5: none ÷ weak

>0.5: it seems data is correlated

>0.8: strong

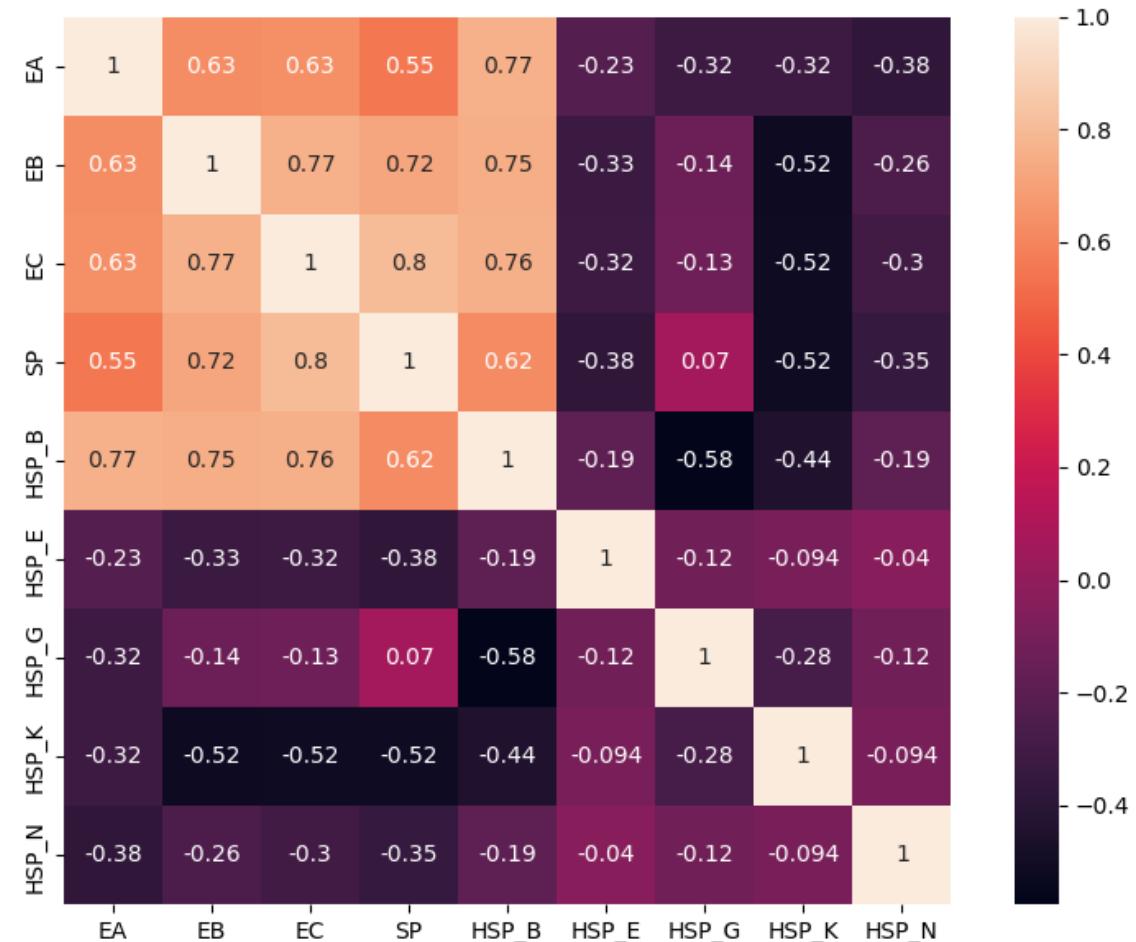
# Looking at correlations

```
import seaborn as sns  
  
corr = df.corr()  
plt.figure(figsize = (8, 6))  
sns.heatmap(corr, square=True, annot=True)  
plt.tight_layout()
```

Some features are colinear:

- EA~HSP\_B
- EB~[EC, HSP\_B]
- EC~HSP\_B

Colinear features are candidates for removal



# Regression plots

```
columns = ['EA', 'EB', 'EC', 'HSP_B', 'HSP_K', 'HSP_E']

fig, axs = plt.subplots(ncols=3, nrows=2, figsize=(10, 6))

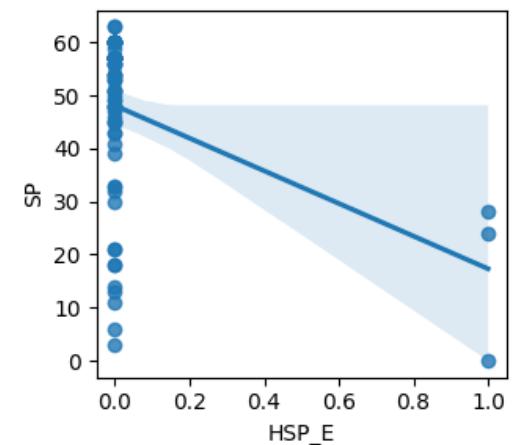
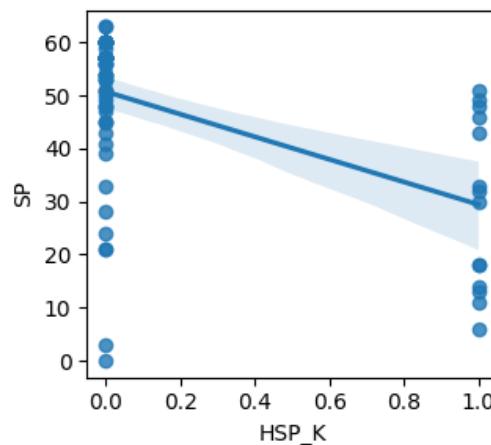
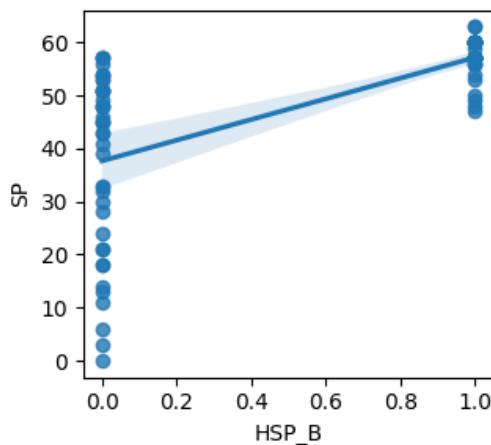
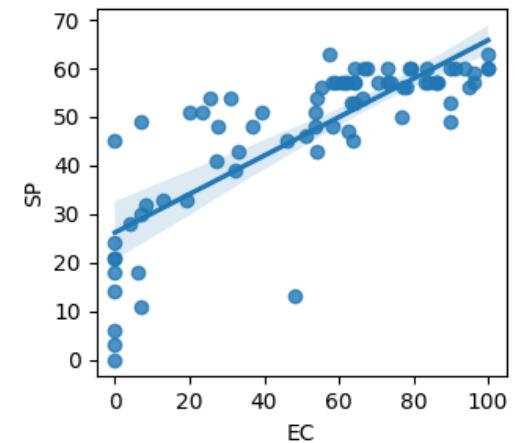
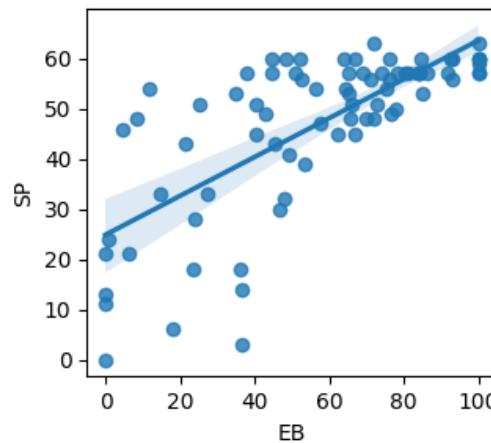
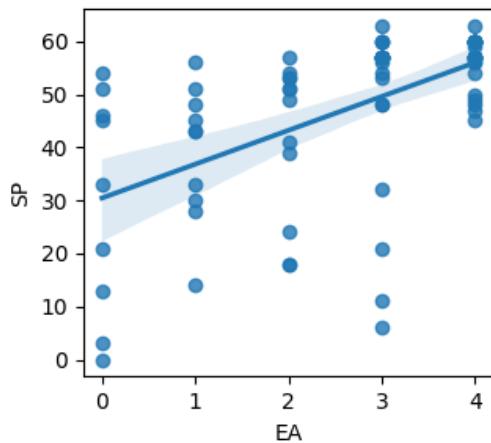
index = 0
axs = axs.flatten()

for i, (k, v) in enumerate(df[columns].items()):

    sns.regplot(y=df['SP'], x=v, ax=axs[i])

plt.tight_layout()
```

# Regression plots

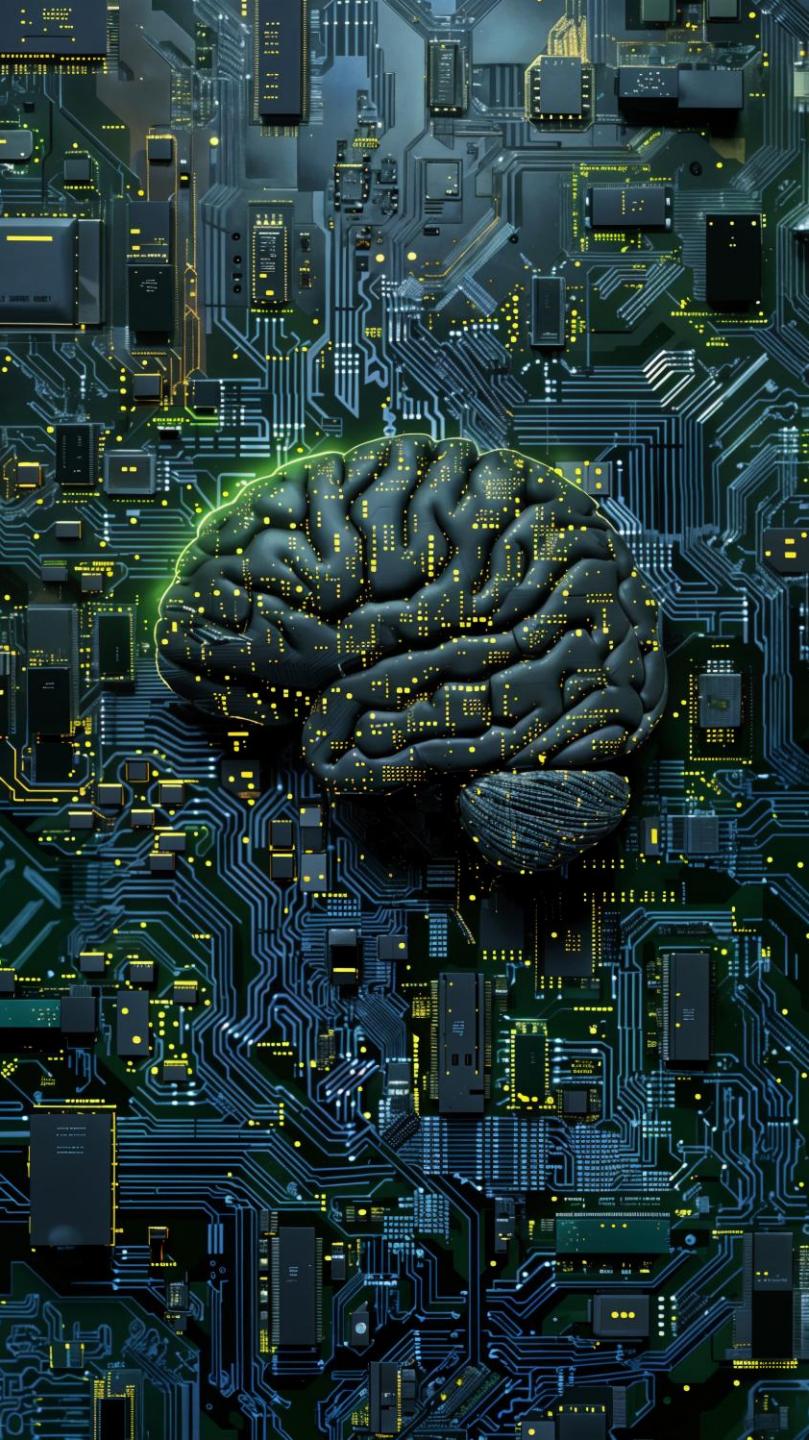


Work, work, work

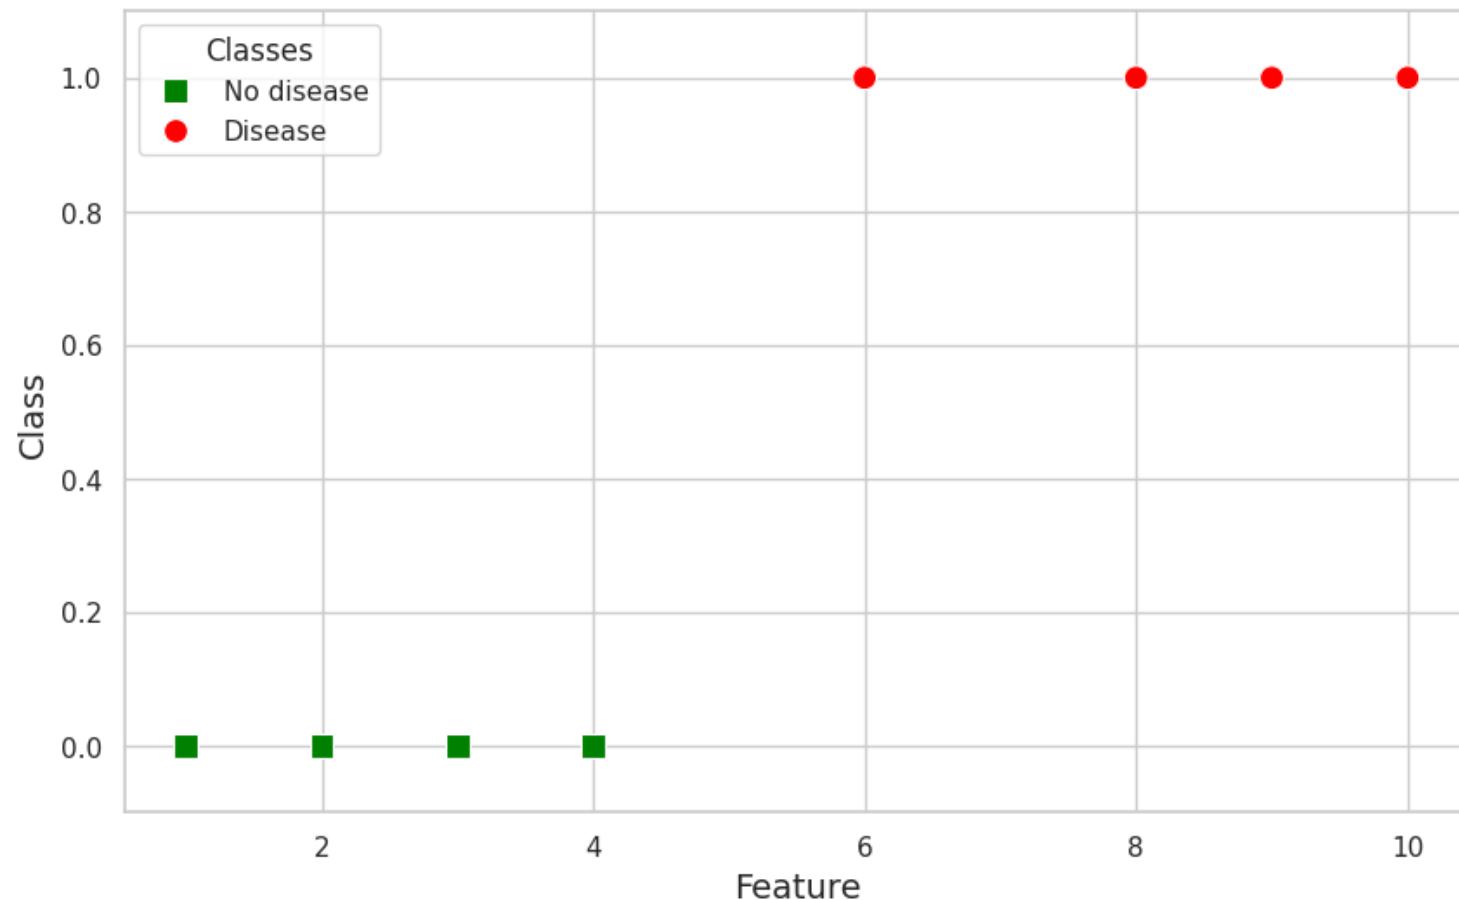


# Classification

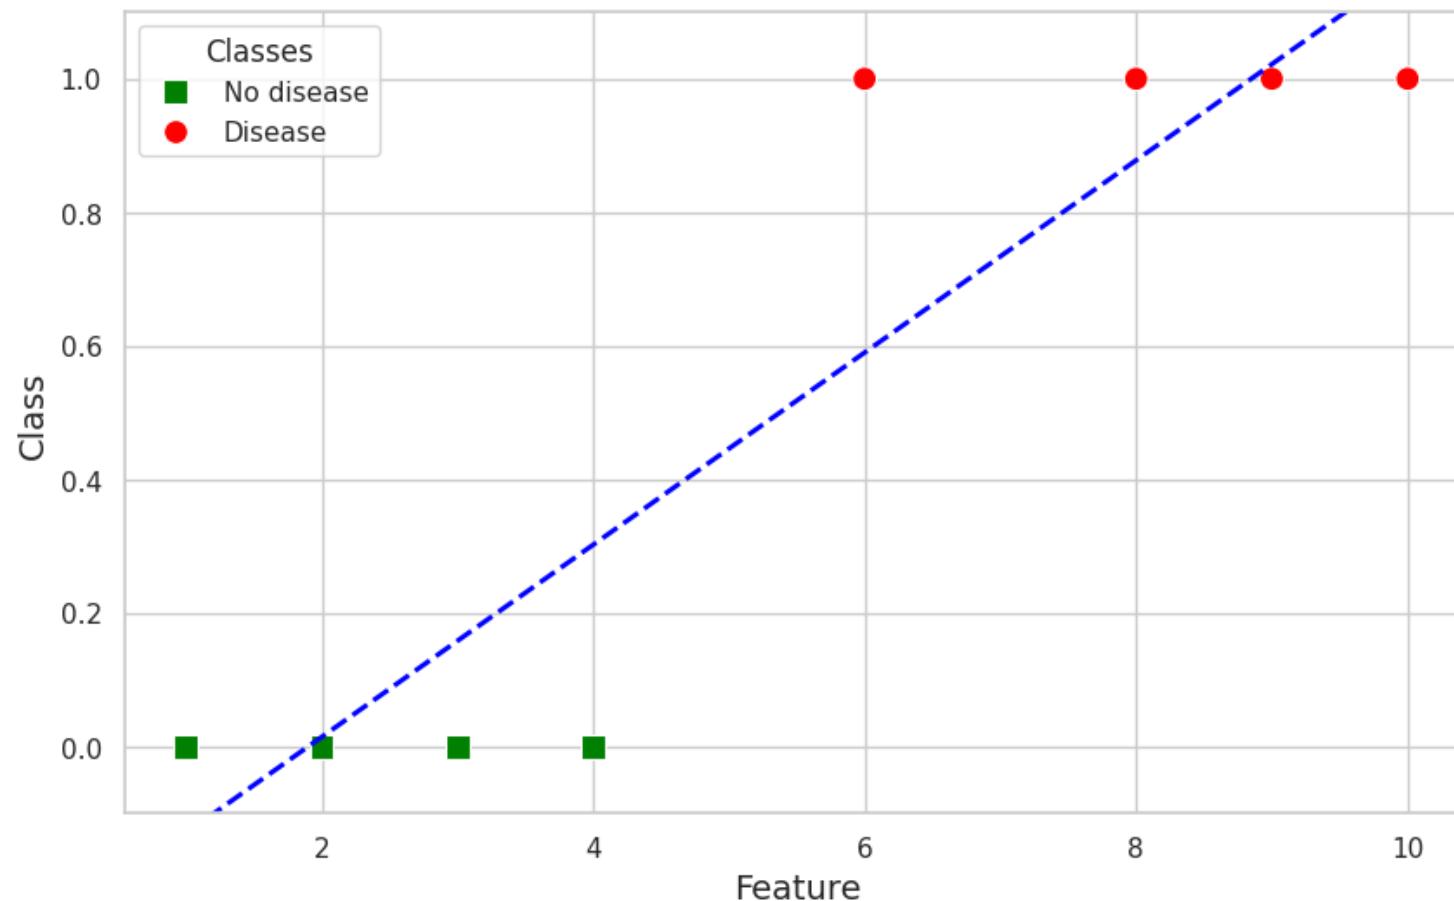
Logistic regression



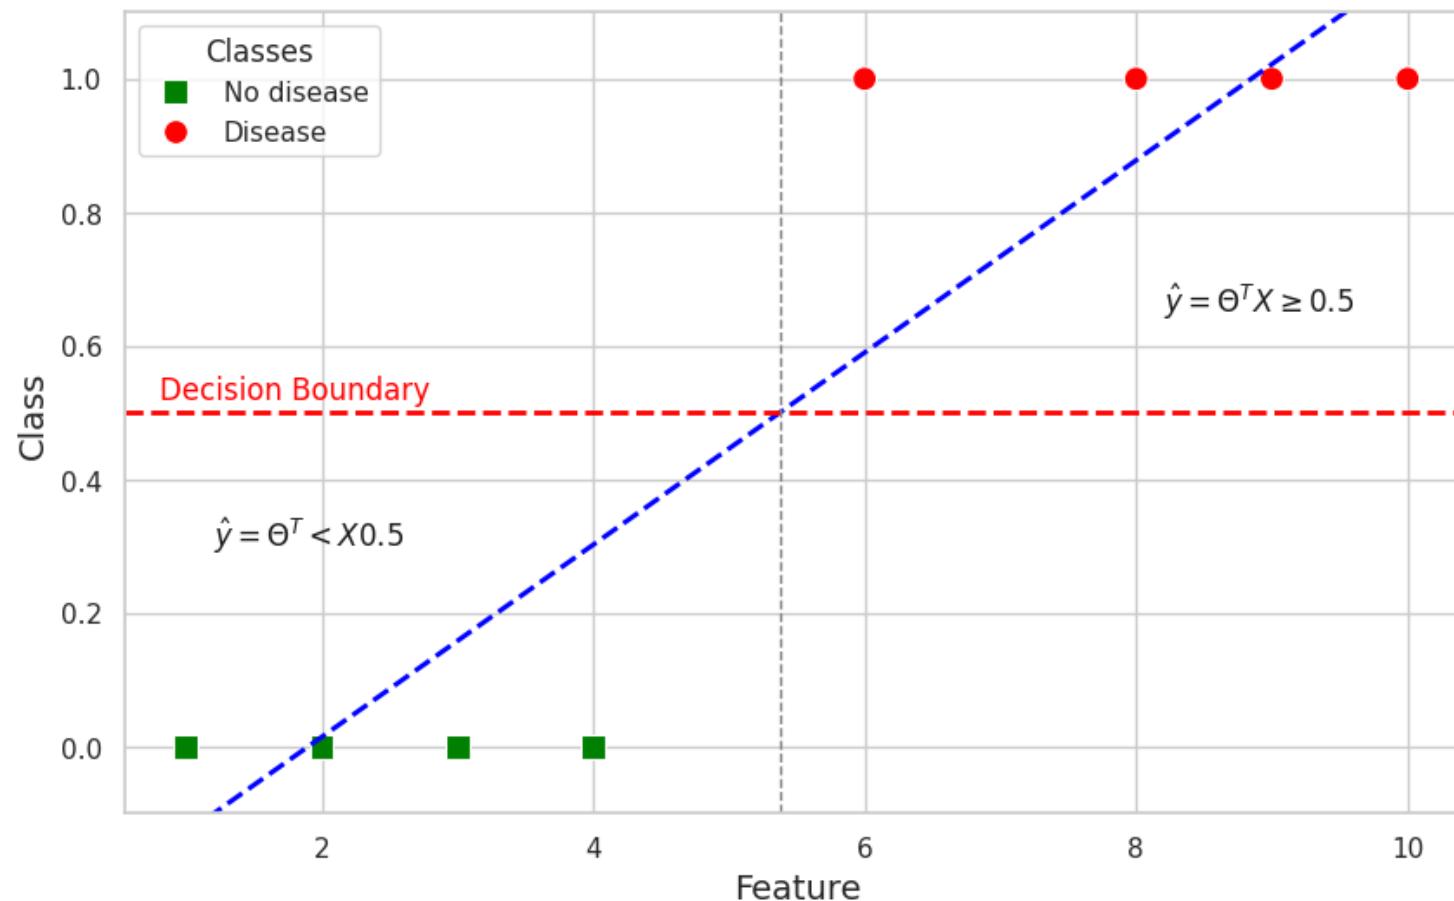
# Simple regression problem



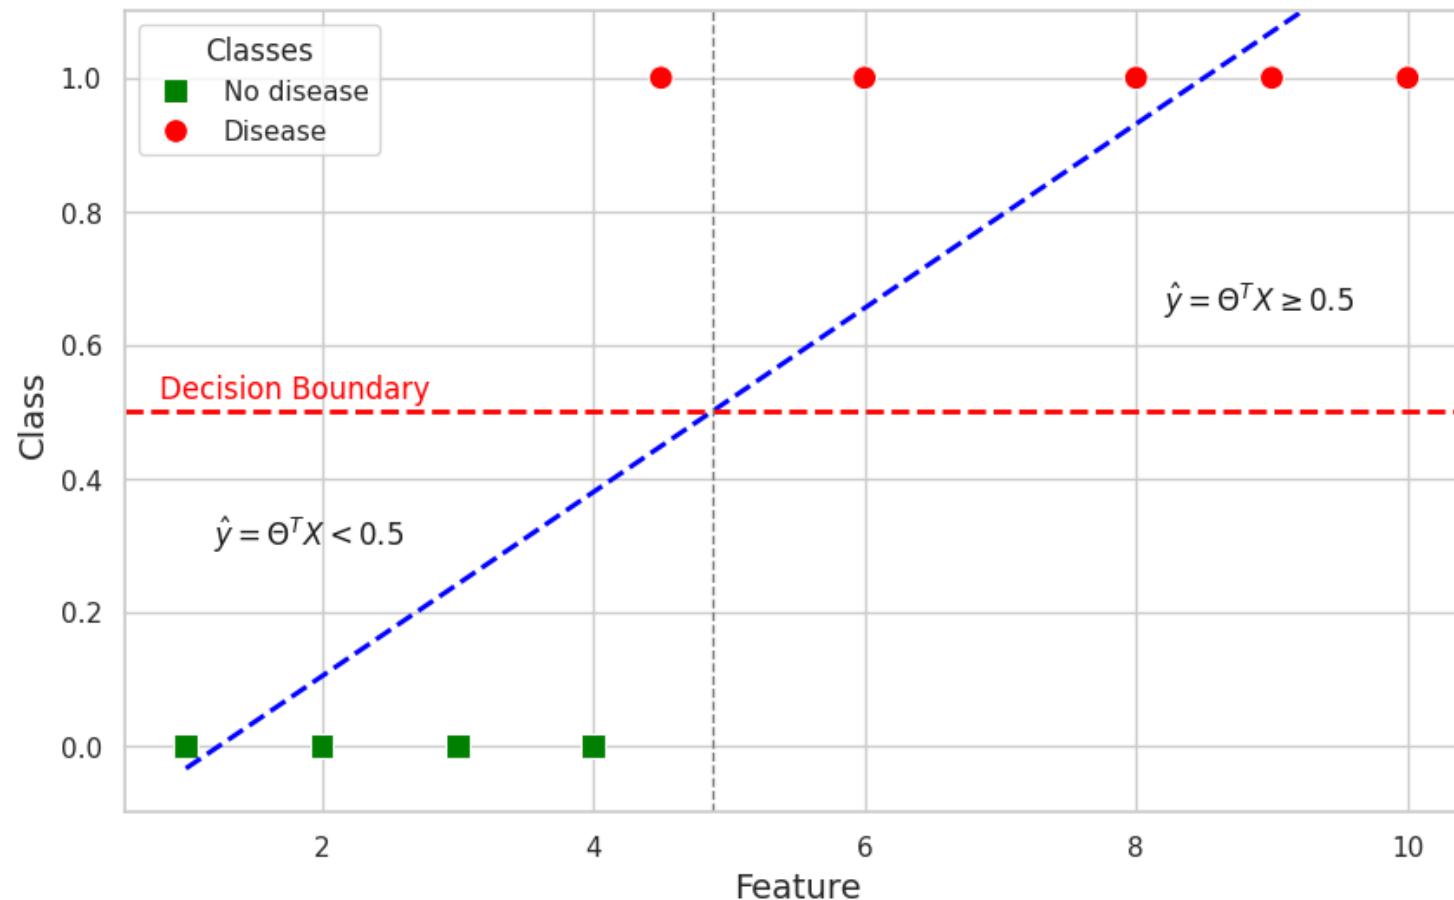
# Simple regression problem



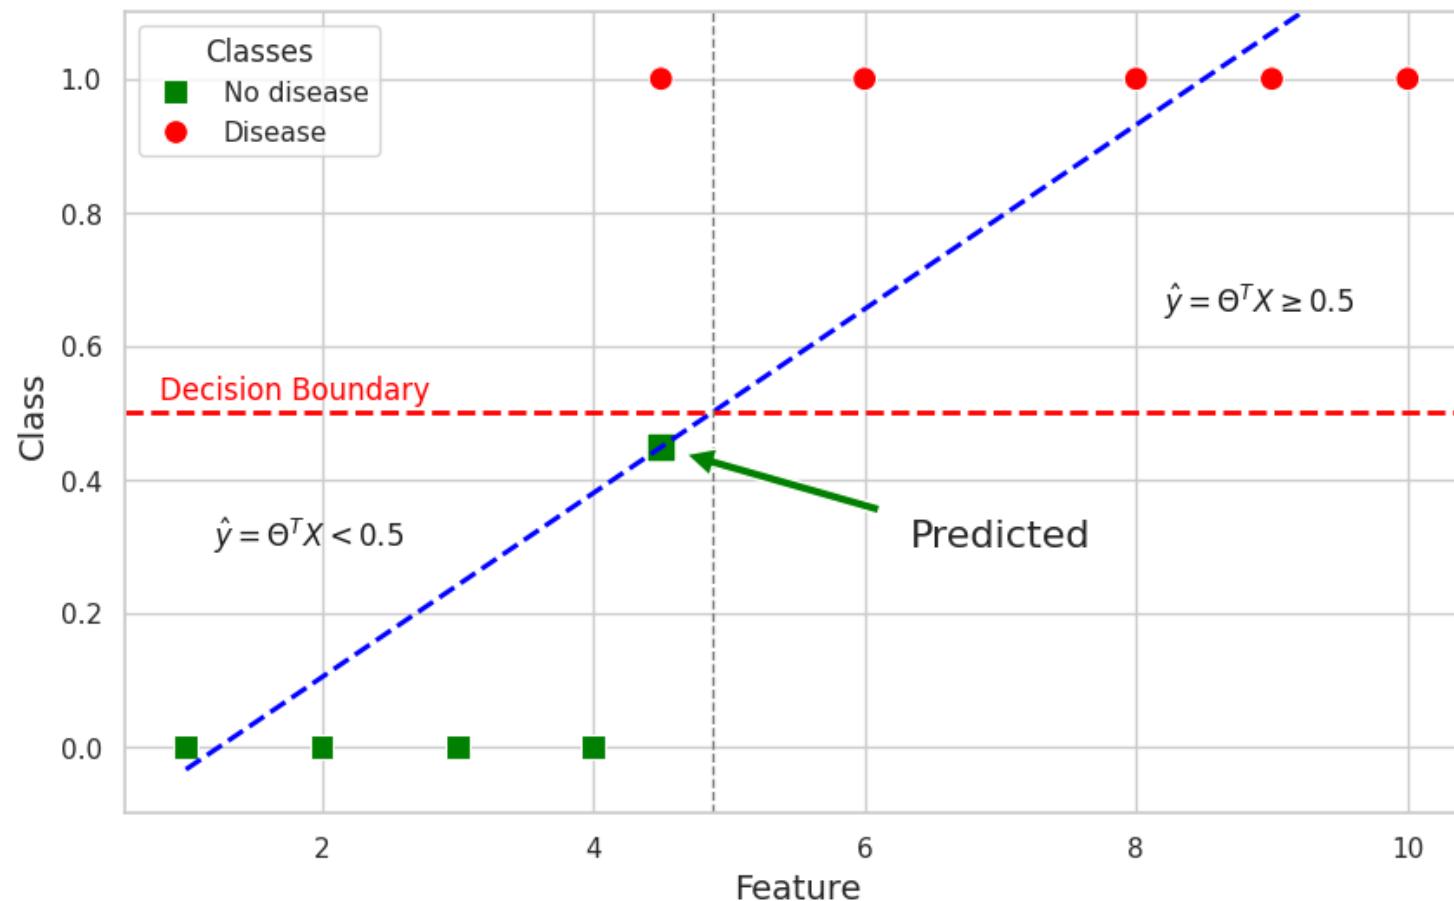
# Simple regression problem



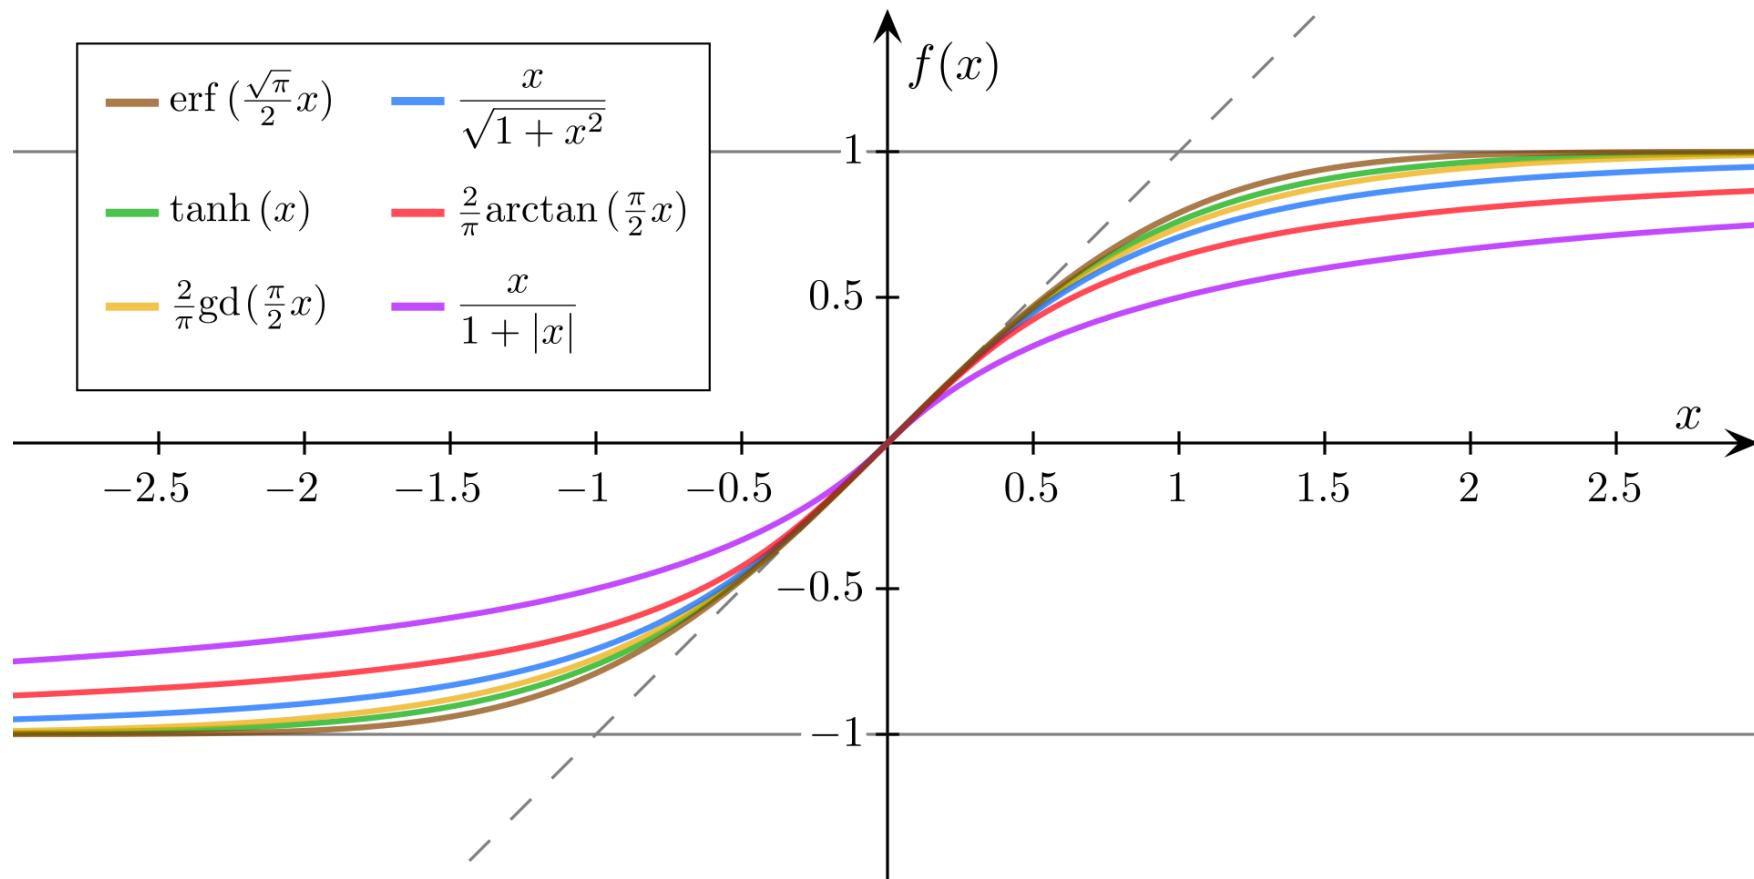
# Simple regression problem



# Simple regression problem



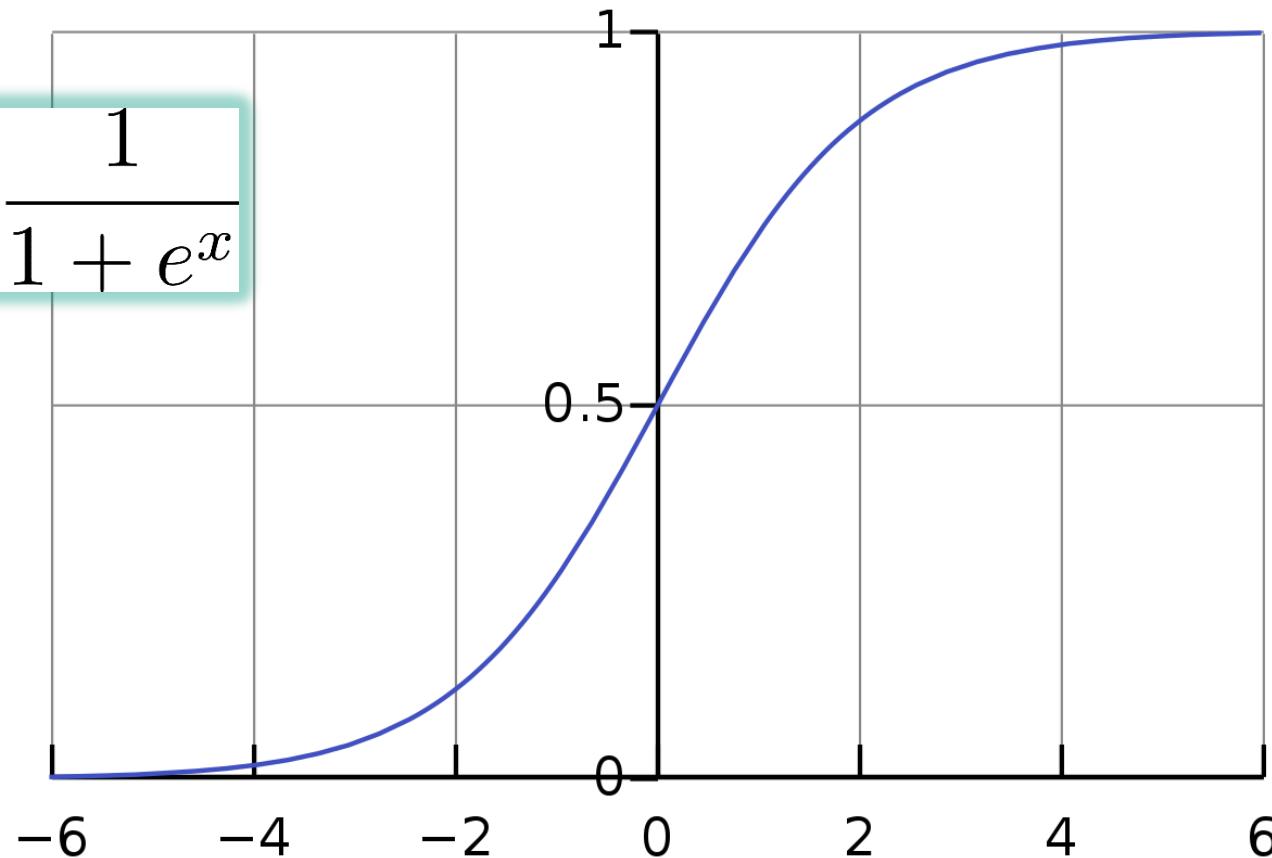
# Welcome the sigmoid



By Georg-Johann (adapted from Geek3), CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=11498624>

# Logistic function

$$f(x) = \frac{1}{1 + e^{-x}}$$



By Qef - Created from scratch with gnuplot, Public Domain, <https://commons.wikimedia.org/w/index.php?curid=4310325>

## Logistic regression

- Linear Regression:

$$h_{\Theta}(X) = \Theta^T X$$

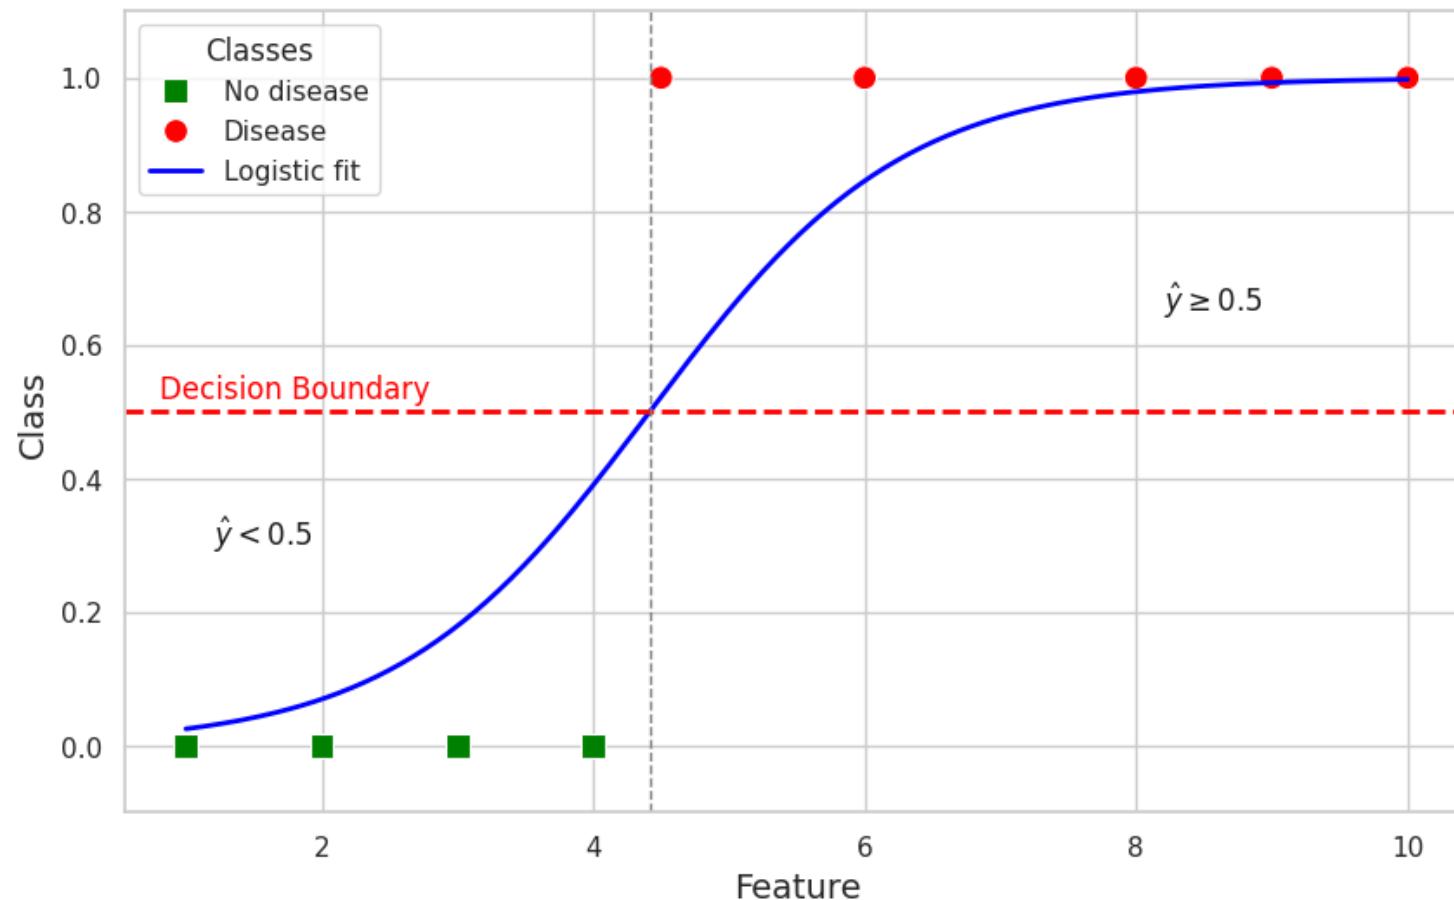
- Logistic Function:

$$f(x) = \frac{1}{1 + e^{-x}}$$

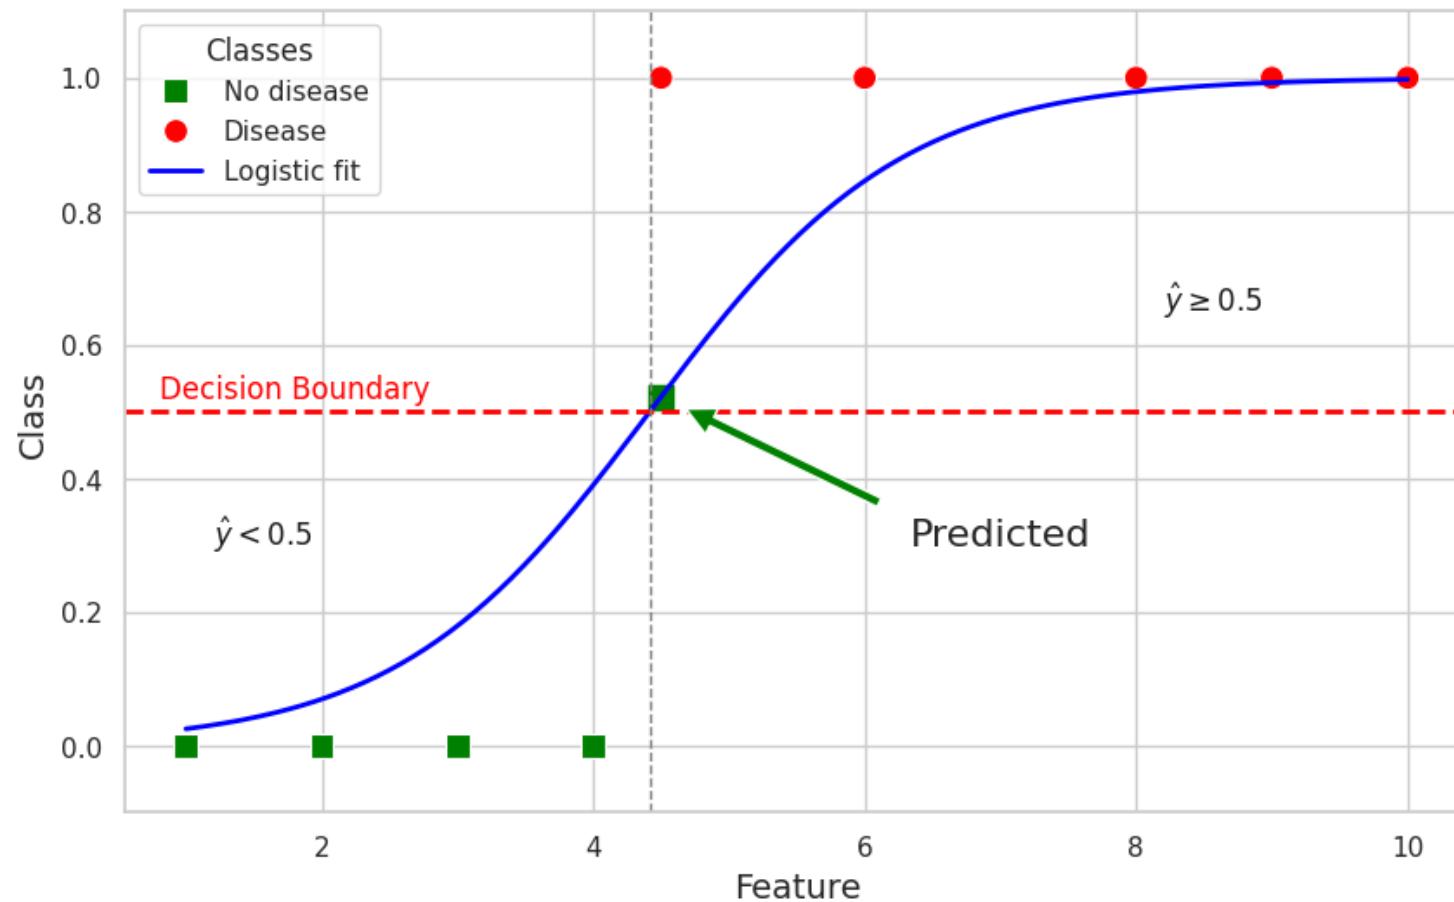
- Logistic Regression:

$$h_{\Theta}(X) = \frac{1}{1 + e^{\Theta^T X}}$$

# Simple regression problem



# Simple regression problem



# Logistic Regression

University of California, Irvine, *Heart Disease Data Set*

```
data = read_csv("heart.csv")
data.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

# Logistic Regression

University of California, Irvine, *Heart Disease Data Set*

`data.describe()`

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
count	303.00	303.00	303.00	303.00	303.00	303.00	303.00	303.00	303.00	303.00	303.00	303.00	303.00	303.00
mean	54.37	0.68	0.97	131.62	246.26	0.15	0.53	149.65	0.33	1.04	1.40	0.73	2.31	0.54
std	9.08	0.47	1.03	17.54	51.83	0.36	0.53	22.91	0.47	1.16	0.62	1.02	0.61	0.50
min	29.00	0.00	0.00	94.00	126.00	0.00	0.00	71.00	0.00	0.00	0.00	0.00	0.00	0.00
25%	47.50	0.00	0.00	120.00	211.00	0.00	0.00	133.50	0.00	0.00	1.00	0.00	2.00	0.00
50%	55.00	1.00	1.00	130.00	240.00	0.00	1.00	153.00	0.00	0.80	1.00	0.00	2.00	1.00
75%	61.00	1.00	2.00	140.00	274.50	0.00	1.00	166.00	1.00	1.60	2.00	1.00	3.00	1.00
max	77.00	1.00	3.00	200.00	564.00	1.00	2.00	202.00	1.00	6.20	2.00	4.00	3.00	1.00

# Logistic Regression

University of California, Irvine, *Heart Disease Data Set*

```
data.corr(method="pearson")['target']
```

```
age      -0.23
sex     -0.28
cp       0.43
thalach   0.42
exang    -0.44
oldpeak  -0.43
slope     0.35
ca       -0.39
thal     -0.34
target    1.00
Name: target, dtype: float64
```

# Logistic Regression

```
from sklearn.linear_model import LogisticRegression

X = data[['thalach']]
y = data['target']

log_reg = LogisticRegression(solver="lbfgs", max_iter=100)

log_reg.fit(X, y)
```

## Class probabilities

```
from sklearn.model_selection import cross_val_score
```

```
cross_val_score(log_reg, X, y, cv=5, scoring="accuracy").mean()
```

```
>> 0.7030054644808743
```

```
log_reg.predict_proba(X)
```

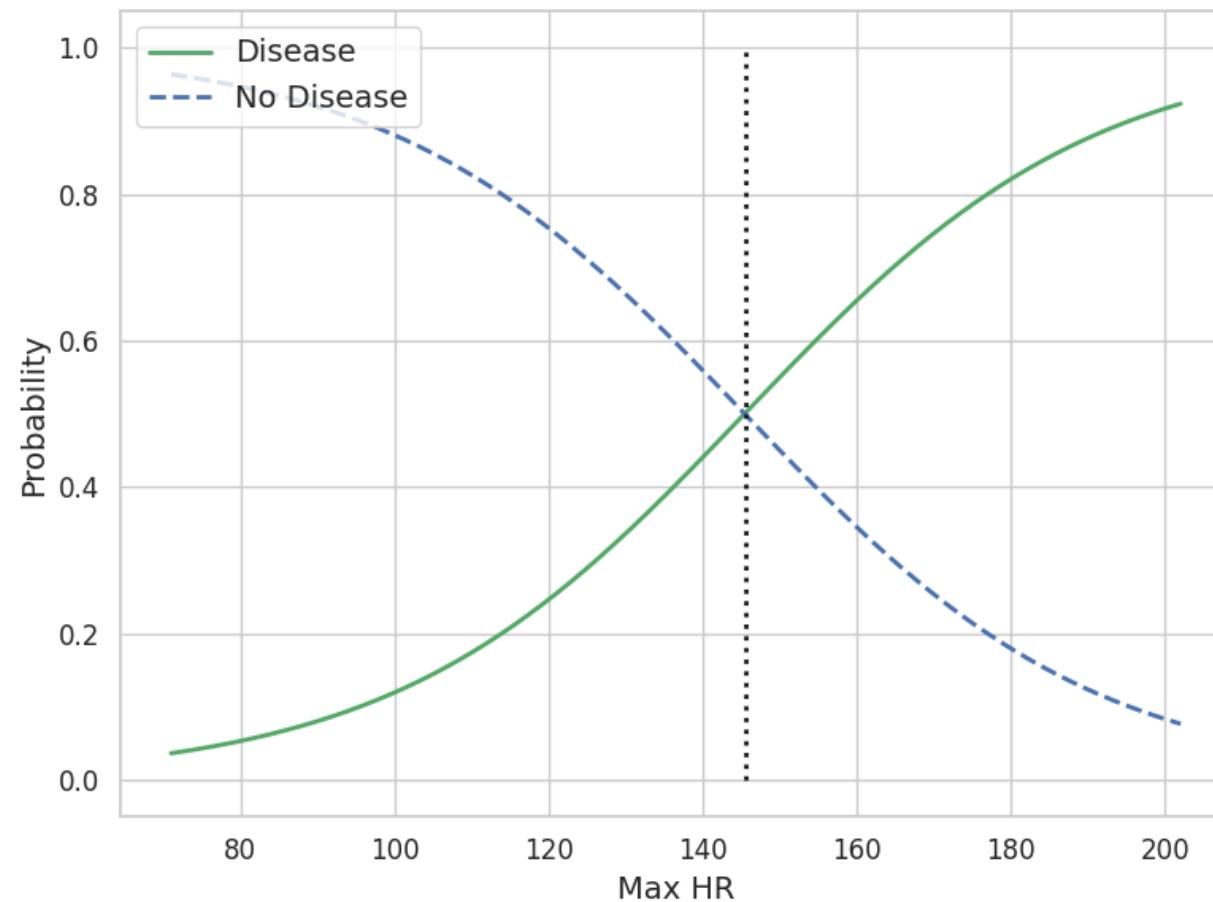


$$h_{\Theta}(X) = \frac{1}{1 + e^{\theta_1 x_1 + \theta_0}}$$



X	0	1
0	0.45	0.55
50	0.46	0.54
...		
100	0.19	0.81
...		
150	0.58	0.42
...		

# One variable regression



## Two variables regression

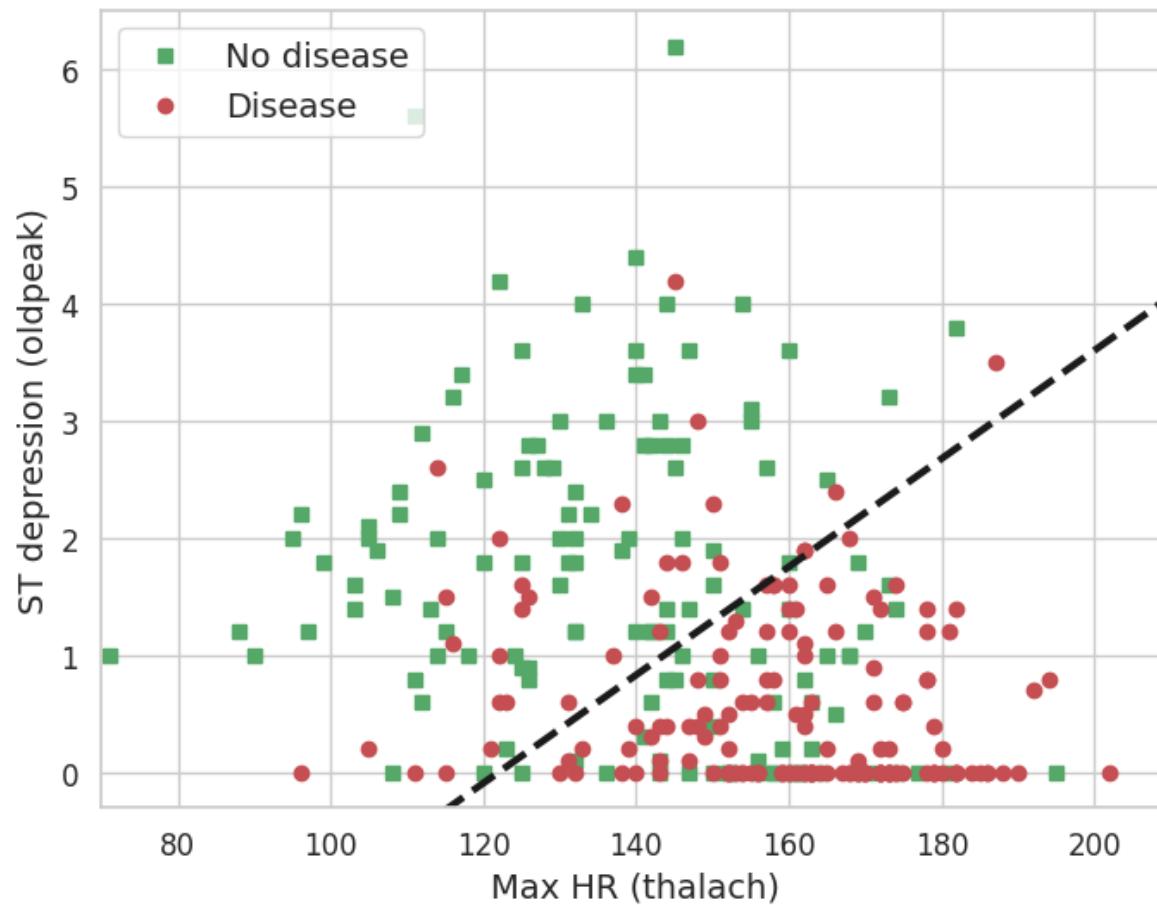
```
X = data[['thalach', 'oldpeak']]  
y = data['target']
```

```
log_reg = LogisticRegression(solver="lbfgs", C=10, max_iter=1000)  
log_reg.fit(X, y)
```

```
>> 0.7325683060109289
```

Plot predicted classes vs. the variables  
Plot decision boundary

## Two variables regression



# Confusion Matrix

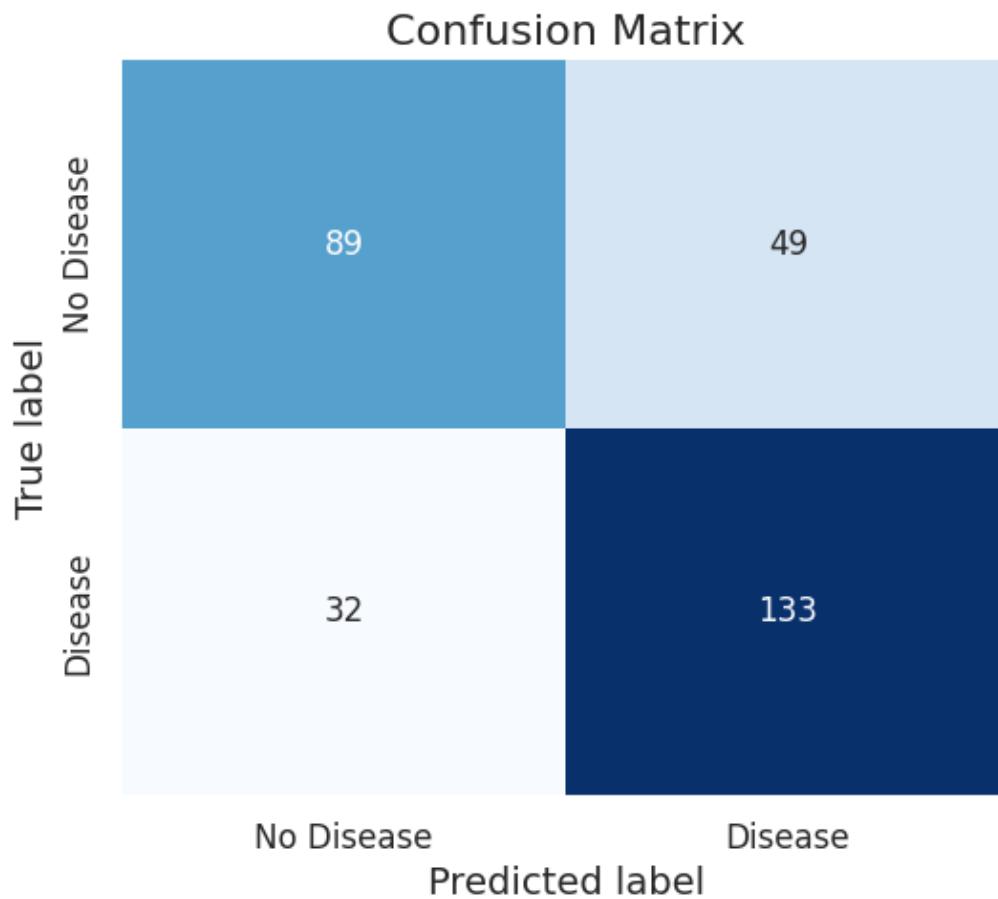
```
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import confusion_matrix

y_pred = cross_val_predict(log_reg, X, y, cv=5)

confusion_matrix(y, y_pred)

>> array([[ 89,  49],
       [ 32, 133]])
```

# Confusion Matrix



# Confusion Matrix

		PREDICTED LABELS	
		Negative	Positive
TRUE LABELS	Negative	True Negative	False Positive
	Positive	False Negative	True Positive

*precision*

*recall*

$$precision = \frac{TP}{TP + FP}$$

$$recall = \frac{TP}{TP + FN}$$

## Confusion Matrix

		PREDICTED	
		Negative	Positive
TRUE	Negative	89	49
	Positive	32	133

```
from sklearn.metrics import precision_score, recall_score
```

```
precision_score(y, y_pred)  
recall_score(y, y_pred)
```

```
>> 0.73
```

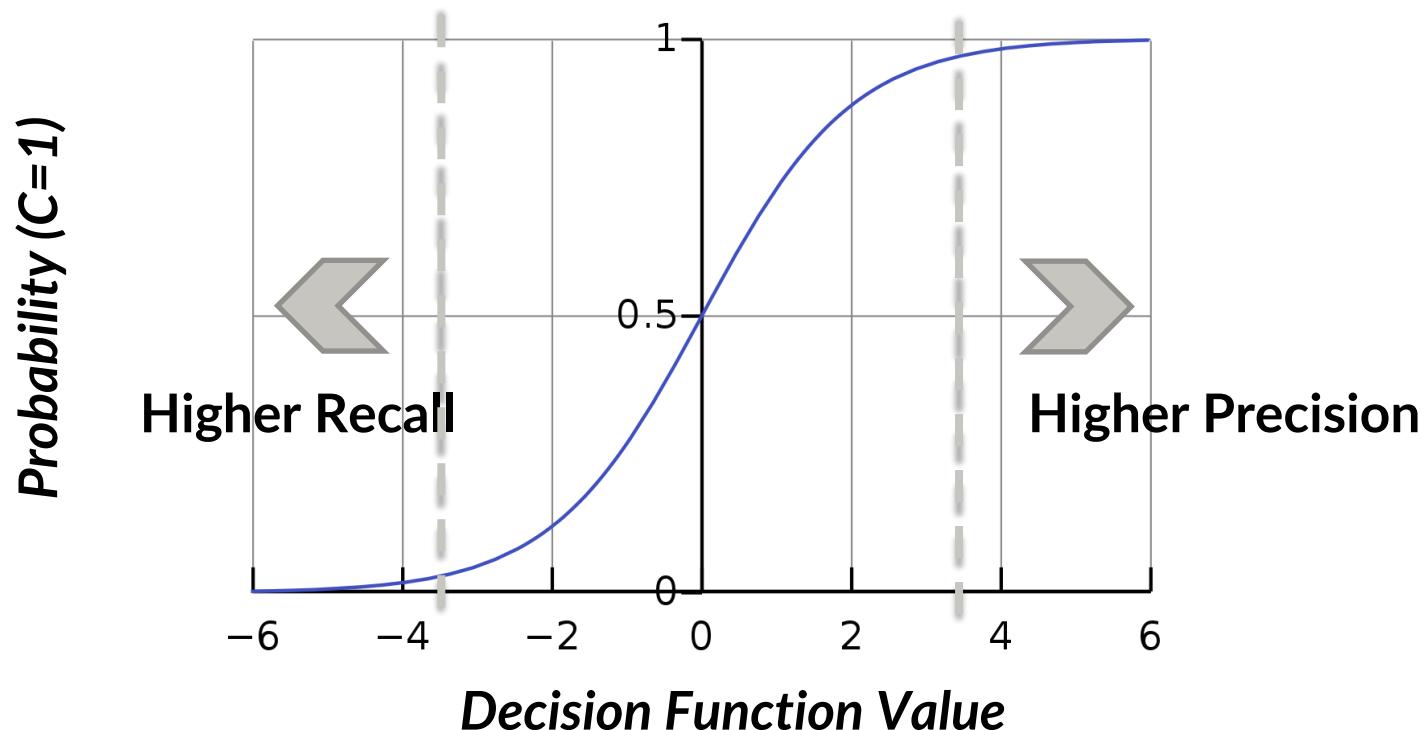
```
>> 0.81
```

# Decision Function

Probability of Positive class:

$$h_{\Theta}(X) = \frac{1}{1 + e^{\Theta^T X}}$$

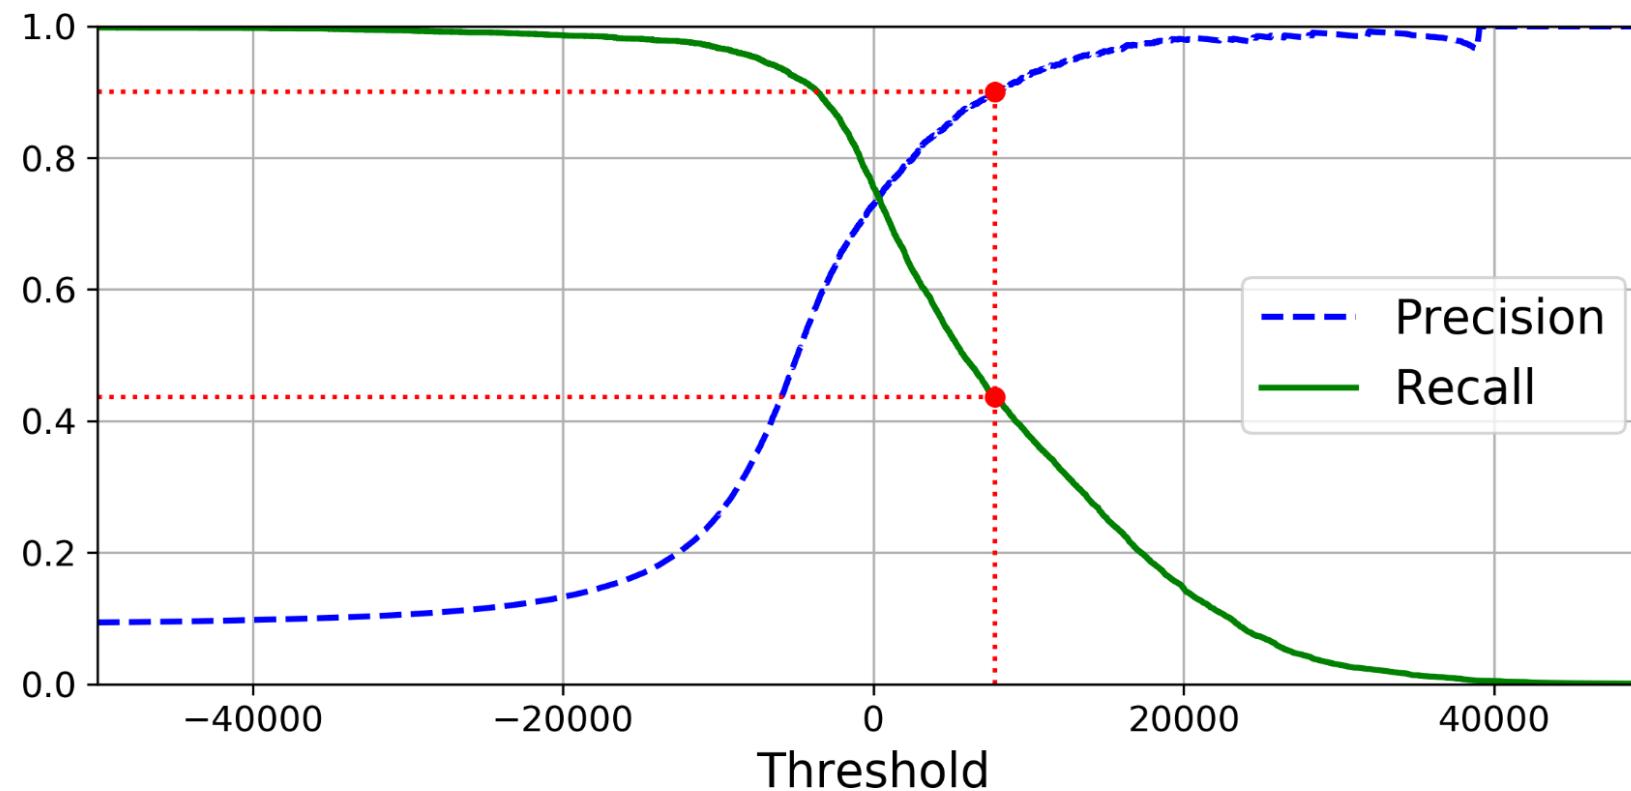
*Decision Function*



## Decision Function Threshold

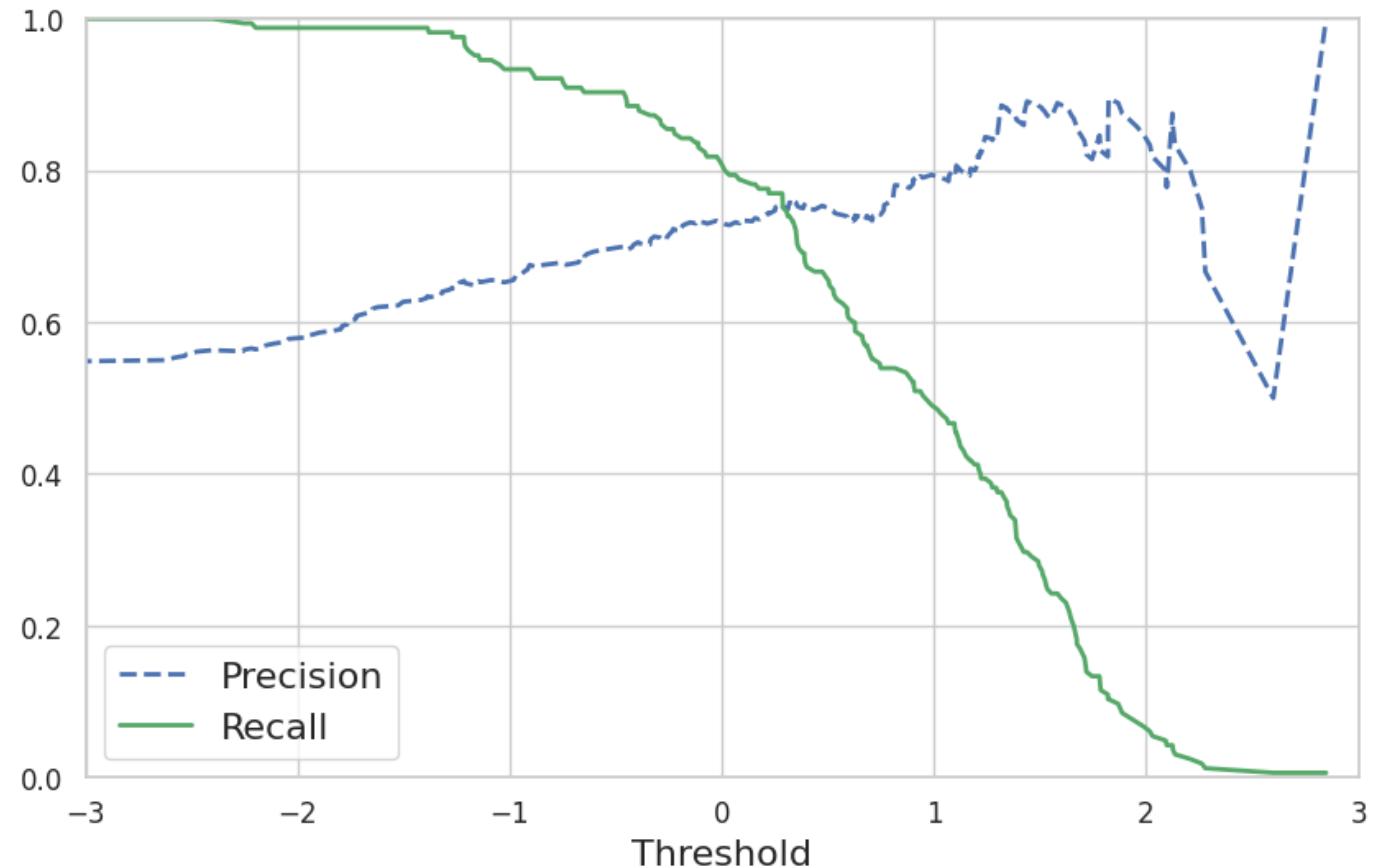
```
y_scores = cross_val_predict(log_reg, X, y, cv=5, method="decision_function")  
  
from sklearn.metrics import precision_recall_curve  
  
precisions, recalls, thresholds = precision_recall_curve(y, y_scores)  
  
# + some more code to print the precisions vs. thresholds  
# and recalls vs. thresholds
```

## Precision vs. Recall



That's the ideal case...  
For our heart disease dataset expect...

# Precision vs. Recall

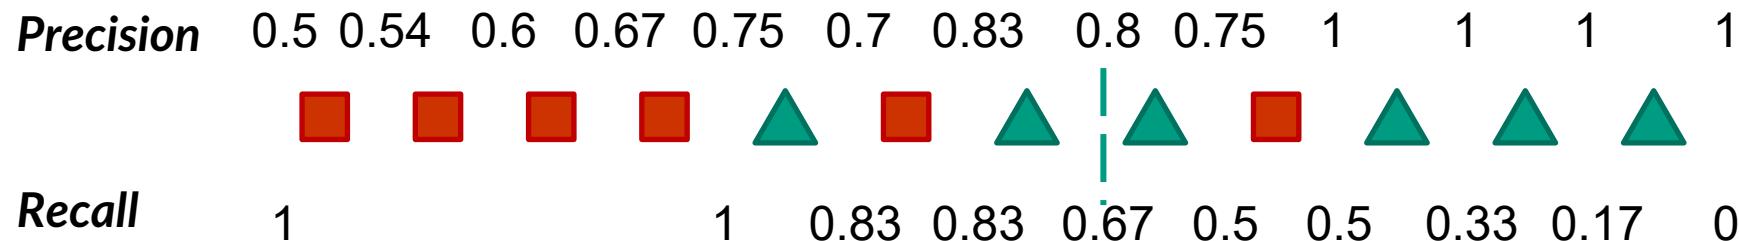


$$Precision = \frac{\text{True Positives}}{\text{Classified as Positives}}$$

$$Recall = \frac{\text{True Positives}}{\text{All Positives}}$$

# Precision vs. Recall

$$Precision = \frac{\text{Positives Classified as Positives}}{\text{Classified as Positives}}$$



$$Recall = \frac{\text{Positives Classified as Positives}}{\text{All Positives}}$$

But there is more! TPR and FPR

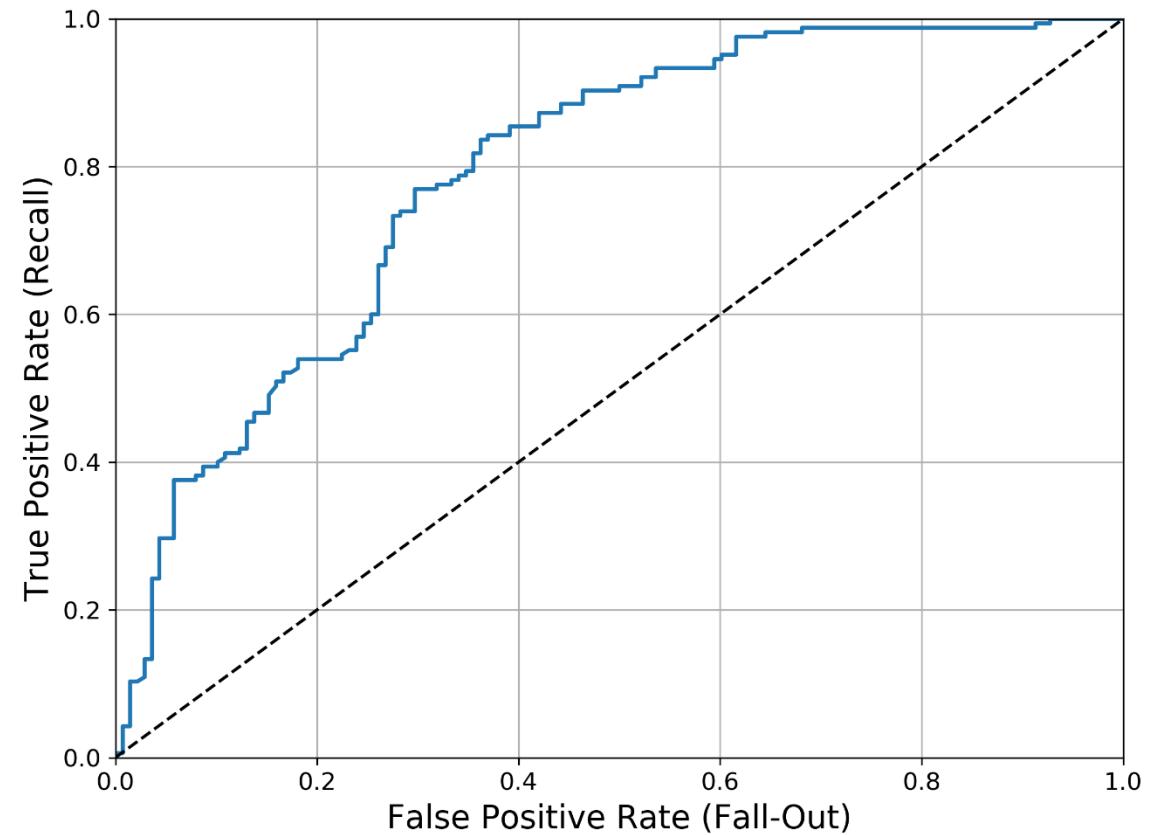
$$\text{recall} = \text{TRP} = \frac{TP}{TP + FN} = \frac{\textit{Positives Classified as Positives}}{\textit{All Positives}}$$

$$\text{fallout} = \text{FPR} = \frac{FP}{TN + FP} = \frac{\textit{Misclassified Nagatives}}{\textit{All Negatives}}$$

\*True Positive Rate, False Positive Rate

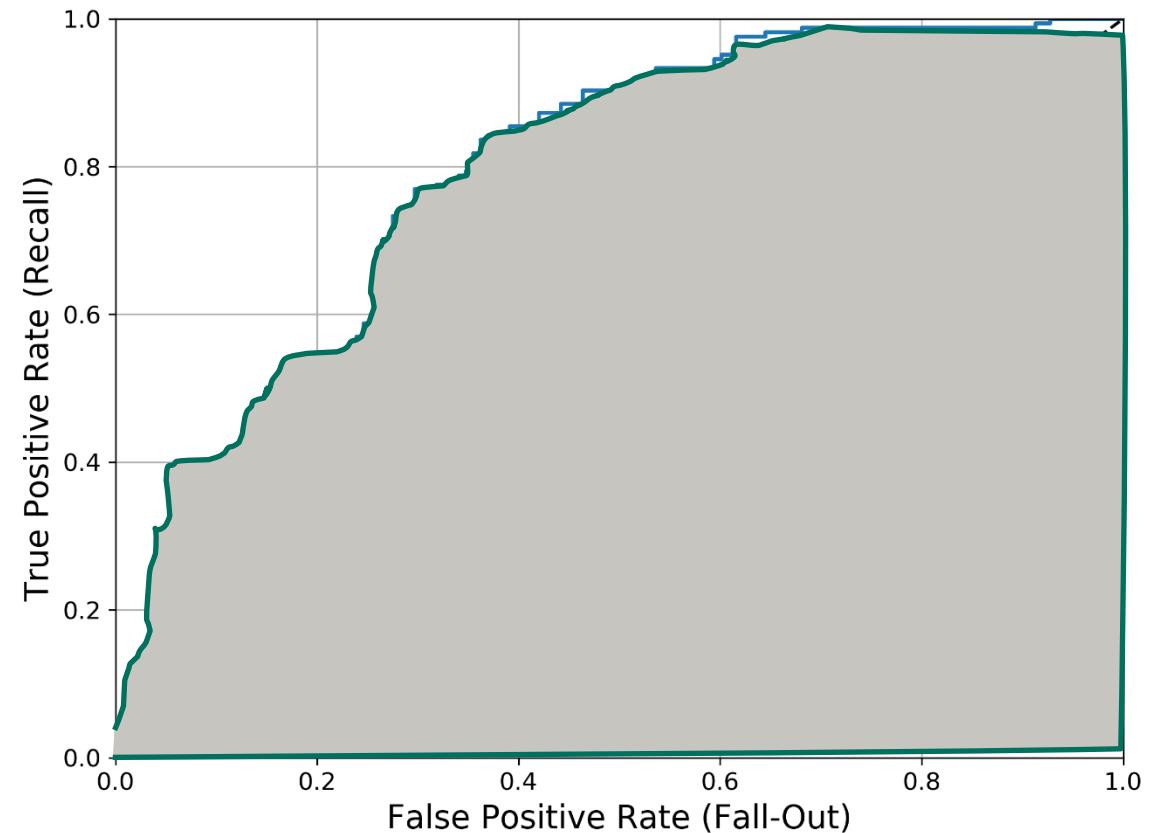
## ROC curve

```
from sklearn.metrics import roc_curve  
  
fpr, tpr, thres = roc_curve(y, y_scores)
```



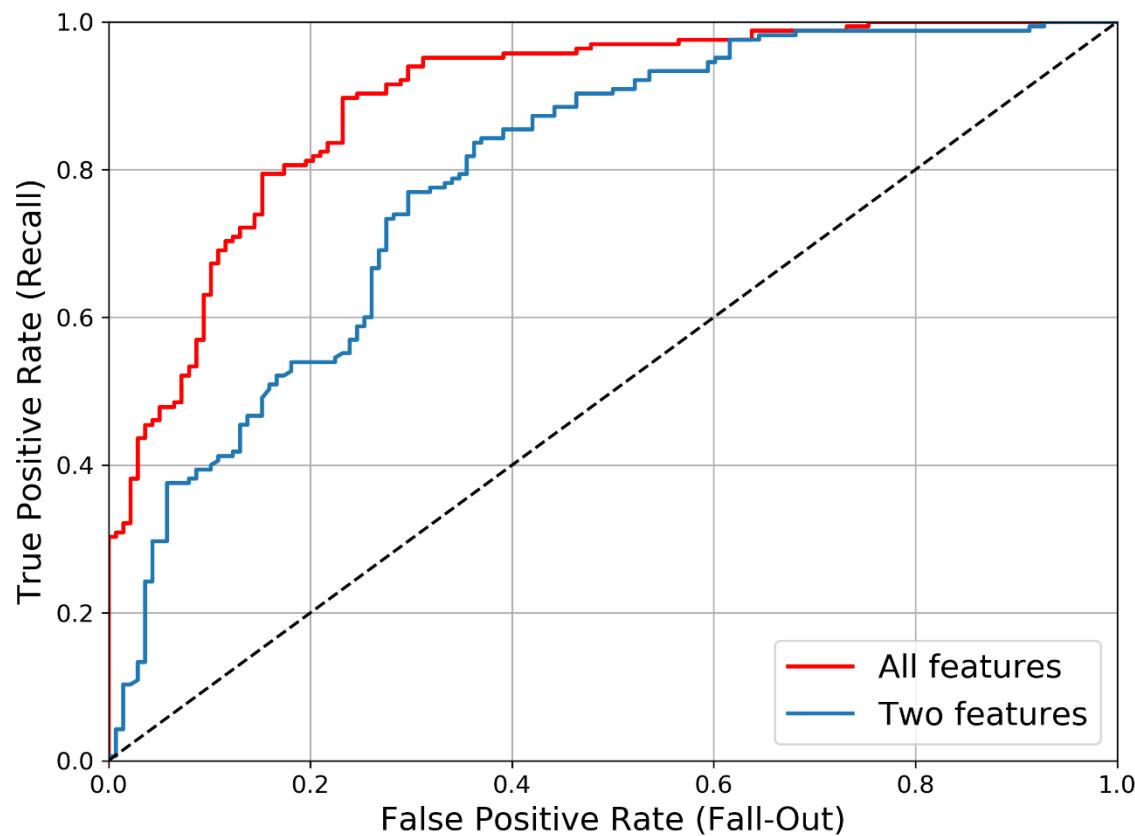
## ROC curve

```
from sklearn.metrics import roc_auc_score  
  
roc_auc_score(y, y_scores)  
  
>> 0.7911725955204216
```



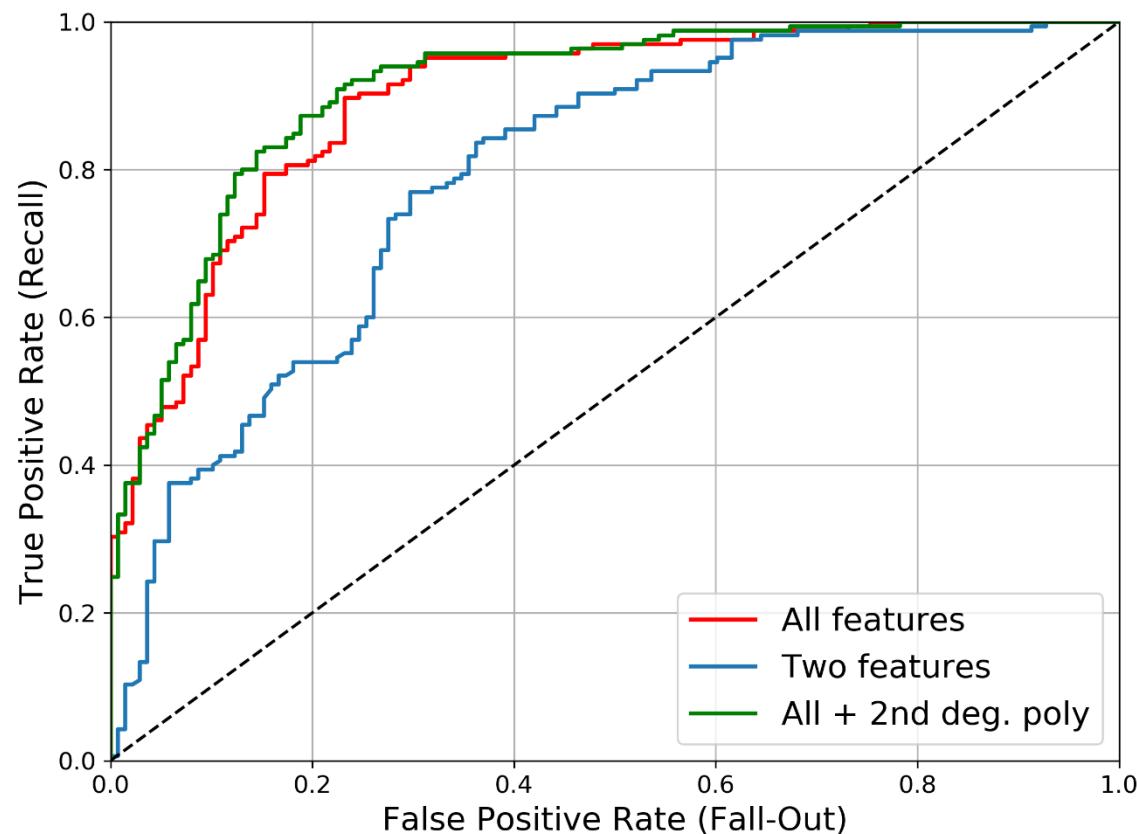
# ROC curve

Logistic Regression: two features vs. all features



# ROC curve

Logistic Regression: pushing harder



Work, work, work

